

Bootstrap Methods

Jonathan Navarrete

June 11, 2017

Introduction

Bootstrap methods in simple terms are methods of *resampling* observed data to estimate the empirical CDF from which the observed data is supposed to have originate from. Suppose we observe *independent* samples x_1, \dots, x_n from pdf/pmf f , and whose CDF F is unobservable (directly). Well, given that $X = (x_1, \dots, x_n)^T$ originates from F , we can use X to generate F_n which is itself an estimate of F . If we sample (with replacement) another set of n observations from F_n , we will have $X^* = (x_1^*, \dots, x_n^*)^T$. This new sample X^* can then generate a CDF, F_n^* which is another estimate of F_n . That is, F_n^* is a bootstrap estimator of F . We can continue this process of resampling with replacement to obtain samples $X_1^*, X_2^*, \dots, X_B^*$ and $F_{n,1}^*, F_{n,2}^*, \dots, F_{n,B}^*$.

In addition to estimating F , there may be a statistic of interest θ (e.g. mean). We can use bootstrap methods to calculate an empirical distribution of θ . From our original sample X we can calculate estimate $\hat{\theta}$. Our bootstrap sample can also be used to calculate an estimate, $\hat{\theta}_1^*, \dots, \hat{\theta}_B^*$.

A simple bootstrap algorithm for *independent* samples X is:

To generate B bootstrap samples, for b in $1, \dots, B$ do


1. Sample x_1, \dots, x_n with replacement; each sample has a probability of $1/n$ of being in the new sample.
2. Calculate $\hat{\theta}_b^*$

We will then observe the empirical distribution of $\hat{\theta}$, $F_{\hat{\theta}}$.

We will use the `mtcars` data set to illustrate a simple implementation.

```
data("mtcars")
mpg = mtcars$mpg
n = length(mpg)
hist(x = mpg, probability = TRUE, xlab = "MPG", main = "Histogram of MPG")
```


Histogram of MPG



```
B = 1000
results = numeric(B)
for(b in 1:B){
  i = sample(x = 1:n, size = n, replace = TRUE)
  bootSample = mpg[i]
  thetaHat = mean(bootSample)
  results[b] = thetaHat
}

hist(x = results, probability = TRUE,
     main = "Bootstrapped Samples of Mean_mpg",
     xlab = "theta estimates")
```

Bootstrapped Samples of Mean_mpg



```
print(table(i)/n)
```

```
## i
##      3      4      6      8     10     11     12     13     15
## 0.03125 0.03125 0.03125 0.03125 0.06250 0.03125 0.03125 0.06250 0.03125
##      16     17     18     20     23     24     26     27     28
## 0.06250 0.03125 0.09375 0.03125 0.09375 0.03125 0.03125 0.03125 0.03125
##      29     30     31     32
## 0.09375 0.06250 0.03125 0.03125
```

As a precaution and note on proper use of bootstrap methods, before embarking on resampling we must ask what variables are *iid* in order to determine a correct bootstrapping approach. Bootstrap methods are *not* a method of generating new data for, say, a regression setting when observed samples are low. In the above example, it is assumed that each observation in the `mpg` data set is independent and identically distributed from an unknown distribution f . However, if there were to have existed some autocorrelation structure (as exist in time-series data) then we would need to adjust our resampling methodology. When dealing with time-series data, we will use a method called *block bootstrap*.

Paired Bootstrapping

Let's continue to work with the `mtcars` data set. Say we wanted to make inferences about the linear regression parameters.


```
library(ggplot2)
```

```
##
## Attaching package: 'ggplot2'
##
## The following object is masked _by_ 'GlobalEnv':
##
##      mpg
```


```
mtcars$am <- as.factor(mtcars$am)

fit = lm(formula = mpg ~ wt + am, data = mtcars)

qplot(x = as.factor(am), y = mpg, data = mtcars, geom = "boxplot",
      main = "Boxplot: MPG ~ AM", ylab = "MPG", xlab = "AM",
      colour = am)
```



```
qplot(x = wt, y = mpg, data = mtcars,
      main = "Boxplot: MPG ~ Weight", ylab = "MPG", xlab = "Weight",
      colour = am)
```



```
## see summary of model
summary(fit)
```

```
##
## Call:
## lm(formula = mpg ~ wt + am, data = mtcars)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.5295 -2.3619 -0.1317  1.4025  6.8782
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  37.32155    3.05464  12.218 5.84e-13 ***
## wt          -5.35281    0.78824  -6.791 1.87e-07 ***
## am1         -0.02362    1.54565  -0.015  0.988
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.098 on 29 degrees of freedom
## Multiple R-squared:  0.7528, Adjusted R-squared:  0.7358
## F-statistic: 44.17 on 2 and 29 DF,  p-value: 1.579e-09
## see coefficients
beta_int = coefficients(fit)[1]
beta_wt = coefficients(fit)[2]
beta_am = coefficients(fit)[3]
```

```

n = dim(mtcars)[1]
B = 1000

results = matrix(data = NA, nrow = B, ncol = 3,
                  dimnames = list(NULL, c("Intercept", "wt", "am")))
for(b in 1:B){
  i = sample(x = 1:n, size = n, replace = TRUE)
  temp = mtcars[i,]
  temp_model = lm(formula = mpg ~ wt + am, data = temp)
  coeff = matrix(data = coefficients(temp_model), ncol = 3)
  if(sum(is.na(coeff)) > 0){
    break
  }
  results[b,] = coeff
}

results <- data.frame(results, check.names = FALSE)

head(results)

```

```

##      Intercept      wt      am
## 1  33.04066 -4.479660  2.9186777
## 2  30.31962 -3.563008  2.8551289
## 3  43.76828 -6.686508 -3.4602491
## 4  32.88873 -4.258204  2.7805738
## 5  38.31424 -5.675963 -0.8044974
## 6  38.94775 -5.658217 -1.2402473

```

```
tail(results)
```

```

##      Intercept      wt      am
## 995  33.08961 -4.151006  0.6965764
## 996  34.37839 -4.829728  2.1764771
## 997  36.81017 -5.265154  0.4666968
## 998  41.58655 -6.222601 -1.6638335
## 999  36.85034 -5.461254 -0.5189622
## 1000 27.38908 -2.698150 -0.7485130

```

```

boot_int = results[, "Intercept"]
boot_wt = results[, "wt"]
boot_am = results[, "am"]

```

```

par(mfrow = c(2,2))
hist(boot_int, main = "Bootstrapped Coefficients for Intercept",
     xlab = "Coefficients for Intercept", probability = TRUE)
abline(v = coefficients(fit)[1], col = "black")

```

```


hist(boot_wt, main = "Bootstrapped Coefficients for Weight",
     xlab = "Coefficients for Weight", probability = TRUE)
abline(v = coefficients(fit)[2], col = "blue")

```


```
hist(boot_am, main = "Bootstrapped Coefficients for AM = 1",
```

```
xlab = "Coefficients for Automatic Transmission", probability = TRUE)
abline(v = coefficients(fit)[3], col = "green")
```


Bootstrapped Coefficients for Intercept



Bootstrapped Coefficients for Weight



Bootstrapped Coefficients for AM = 1



Now we can estimate bias

```
bias_int = mean(boot_int - beta_int)
print(bias_int)
```

```
## [1] 0.5395886
```

```
bias_wt = mean(boot_wt - beta_wt)
print(bias_wt)
```

```
## [1] -0.1761782
```

```
bias_am = mean(boot_am - beta_am)
print(bias_am)
```

```
## [1] -0.1108488
```

```
## incorporate our bias into the coefficients
## we now have bias corrected coefficients
```

```
intercept = beta_int - bias_int
print(intercept)
```

```
## (Intercept)
```

```
## 36.78196
```

```
wt = beta_wt - bias_wt
print(wt)
```

```
##          wt
## -5.176633
am = beta_am - bias_am
print(am)

##          am1
## 0.08723355
```

Define Bias as $Bias(\theta) = E[\theta^*] - \theta$, where in our scenario we have $Bias(\hat{\theta}) = E[\hat{\theta}^*] - \hat{\theta}$. Our bootstrap bias corrected estimates are then $\hat{\theta}_{BC} = \hat{\theta} - Bias(\hat{\theta})$.