

Introduction: Monte Carlo Methods

Jonathan Navarrete

July 30, 2017

Introduction to Monte Carlo

Monte Carlo methods are methods for generating random variables directly or indirectly from target distributions. We generate random variables to estimate p-values or parameters.

Applications of monte carlo methods are in hypothesis testing and Bayesian computation.

Simulation: Gambler's ruin

Consider two gamblers, persons A and B, who start to gamble in a zero-sum game with stakes $\$x$ and $\$(b-x)$, respectively. At each round, each gambler puts up a stake of $\$h$. The probability that A wins a round is p , while the probability that B wins a round is $q = 1 - p$. We wish to compute the probability that A ultimately wins the game. Let us define $v(x, t)$ to be the probability that A ultimately wins the game starting with capital $\$x$ on or before the t th round. Similarly, $u(x, t)$ is the probability that B wins the game with their stake of $b - x$ on or before the t th round.

Each of three variables v, u , and w is bounded below by zero and above by 1. Moreover, u and v are nondecreasing in t . w is nonincreasing in t . Thus we can take limits of each of these as t goes to infinity. We shall call these limits $v(x)$, $u(x)$, and $w(x)$, respectively.

Gambler's ruin, pt.2


Gambler's ruin (fallacy) is the belief that a certain event is *more* likely to occur given the past history. In an experiment where there is a coin toss with probability of seeing heads as 0.5. Each flip of a coin has the same probability of landing on heads regardless of what the previous lands were.

Imagine a gambler on a roulette table. Say the gambler starts with \$10. In this game, the gambler “wins” when they earn a total of \$20 (that is they must play the game until they've earned \$10 on top of their starting \$10). For each game, there is a probability of winning, $p = 0.473$. Then, can we see how many turns until he/she wins (or loses)?

Gambler's ruin, pt.3

```
set.seed(678)
N = 200
income = 10
games = 2*(runif(N)<0.473) - 1 ## generate 1s and -1s
out = cumsum(games) + income

plot(1:N, out, type = "l", xlab = "games", ylab = "income")
abline(h = 0, col = "red")
```



Monte Carlo Hypothesis Testing

There are two ways that the Chi-squared test is used:

1. comparing the observed distribution to some theoretical distribution pre-specified ahead of time: to test the *Goodness of fit* of the theoretical distribution to the observations;
2. testing for *independence* between different factors (which, technically, is just a specific theoretical distribution, with some extra parameters that must be estimated from the data).

To review the Chi-squared test, follow the *link*

Monte Carlo Hypothesis Testing, pt. 2

However, a disadvantage of the chi-square test is that it requires a sufficient sample size in order for the chi-square approximation to be valid. When cell counts are low, say, below 5 asymptotic properties do not hold well. Therefore, a simple chi-squared test may report an invalid p-value which would increase a Type I error rate.

Monte Carlo Hypothesis Testing, pt. 3

A solution is to use Monte Carlo simulation to generate samples from the null distribution in order to estimate a more accurate p-value to our hypothesis.

```
study = matrix(data = c(21, 2,  
                        15, 3), nrow = 2, ncol = 2, byrow = TRUE,  
               dimnames = list(c("surgery", "radiation"),  
                               c("controlled", "not controlled"))  
  
print(study)
```

```
##           controlled not controlled  
## surgery           21             2  
## radiation          15             3
```

Set up some functions in order to generate our Chi-squared statistic and Monte Carlo p-value.

Monte Carlo Hypothesis Testing, pt. 4

```
## set up

## function will generate chi-squared statistics
## using the expected distribution of the data
simulateChisq <- function(B, E, sr, sc){
  results = numeric(B)
  for(i in 1:B){
    ## review r2dtable documentation
    dat = unlist(r2dtable(1, sr, sc))
    M = matrix(dat, ncol = length(sc), nrow = length(sr))
    val = sum( sort( (M - E)^2 / E, decreasing = TRUE))
    results[i] = val
  }
  return(results)
}
```

Monte Carlo Hypothesis Testing, pt. 5

```
ChisqTest <- function(data, Simulations){  
  x = data  
  B = Simulations  
  n <- sum(x)  
  sr <- rowSums(x)  
  sc <- colSums(x)  
  E <- outer(sr, sc, "*")/n ## ORDER MATTERS  
  dimnames(E) <- dimnames(study)  
  tmp <- simulateChisq(B, E, sr, sc)  
  Stat <- sum(sort((x - E)^2/E, decreasing = TRUE))  
  pval <- (1 + sum(tmp >= Stat))/(B + 1)  
  rawPVal = pchisq(q = Stat, df = 2, lower.tail = FALSE)  
  out = list(PearsonStat = Stat,  
             MonteCarloPVal = pval,  
             rawPVal = rawPVal)  
  return(out)  
}
```

Monte Carlo Hypothesis Testing, pt. 6

We then generate our test statistics.

```
set.seed(123)

results <- ChisqTest(study, 10000)

print(results)
```

```
## $PearsonStat
## [1] 0.5991546
##
## $MonteCarloPVal
## [1] 0.6417358
##
## $rawPVal
## [1] 0.7411314
```

```
## compare against chisq.test()
```

Though our ultimate decision to support the null hypothesis of dependence is not a surprise, our results show that the Monte Carlo p-value is greater than the raw p-value obtained from the calculated χ^2 statistic indicating more support for the null hypothesis. Readers should compare these results against R's `chisq.test` function.

Bayesian Example

Here is an example taken from *Bayesian Ideas and Data Analysis* by Christensen et al.

$$y|\theta \sim \text{Bin}(2430, \theta) \text{ and } \theta \sim \text{Beta}(12.05, 116.06)$$

This is a beta-binomial problem. There is a beta prior distribution on θ . Beta is conjugate to the binomial distribution, see: Conjugate priors. Bayesian analysis uses prior information combined with observed data to update a probability distribution, posterior distribution, from which we can obtain a probability value. The new probability distribution, posterior, describes knowledge about the unknown parameter θ from historical beliefs (e.g. previous experiments, reports, etc.) and current observed data.

Bayesian Example, pt. 2

The Bayesian data model is then

$$y|\theta \sim \text{Bin}(n, \theta) \text{ and } \theta \sim \text{Beta}(a, b)$$

The resulting posterior distribution is then


$$\theta|y \sim \text{Beta}(y + a, n - y + b)$$

Bayesian Example, pt. 3

We can now simulate the posterior distribution

```
N = 10^4
set.seed(123)
x = rbeta(n = N, shape1 = 219 + 12.05, shape2 = 2430 - 219 + 116.06)
d = density(x)
hist(x = x, probability = TRUE,
     main = "Beta Posterior Distribution",
     xlab = "x (posterior values)", ylab = "Density",
     ylim = c(0,80))
lines(x = d$x , y = d$y, type = "l", col = 2)
```

Beta Posterior Distribution



Bayesian Example, pt. 4

We can tell the VP that the true probability lies between 7.9% and 10.2%, with median probability of 9%.