# From EM to Data Augmentation: The Emergence of MCMC Bayesian Computation in the 1980s

**Martin A. Tanner and Wing H. Wong**

*Abstract.* It was known from Metropolis et al. [*J. Chem. Phys.* **21** (1953) 1087–1092] that one can sample from a distribution by performing Monte Carlo simulation from a Markov chain whose equilibrium distribution is equal to the target distribution. However, it took several decades before the statistical community embraced Markov chain Monte Carlo (MCMC) as a general computational tool in Bayesian inference. The usual reasons that are advanced to explain why statisticians were slow to catch on to the method include lack of computing power and unfamiliarity with the early dynamic Monte Carlo papers in the statistical physics literature. We argue that there was a deeper reason, namely, that the structure of problems in the statistical mechanics and those in the standard statistical literature are different. To make the methods usable in standard Bayesian problems, one had to exploit the power that comes from the introduction of judiciously chosen auxiliary variables and collective moves. This paper examines the development in the critical period 1980–1990, when the ideas of Markov chain simulation from the statistical physics literature and the latent variable formulation in maximum likelihood computation (i.e., EM algorithm) came together to spark the widespread application of MCMC methods in Bayesian computation.

*Key words and phrases:* Data augmentation, EM algorithm, MCMC.

## 1. INTRODUCTION

This paper surveys the historical development of MCMC methodology during a key time period in Bayesian computation. As of the mid-1980s, the Baye-

*Martin A. Tanner is Professor, Department of Statistics, Northwestern University, Evanston, Illinois 60208, USA e-mail: mat132@northwestern.edu. Wing H. Wong is Professor, Departments of Statistics and Health Research and Policy, Stanford University, Stanford, California 94305, USA e-mail: whwong@stanford.edu.*

sian community was focused on Gaussian quadrature type methods, Laplace approximations and variants of importance sampling as the main computational tools in Bayesian analysis. Among more dogmatic Bayesians, the use of Monte Carlo was met with resistance and viewed as antithetical to Bayesian principles. MCMC techniques published in the statistical physics and image analysis literature were seen by the Bayesian computational community as techniques for specialized problems. However, by the early 1990s MCMC-based approaches have become a mainstay in computational Bayesian inference. The purpose of this paper is to review the events that led to this remarkable development. In particular, we examine the critical decade of 1980–1990 when the ideas of Markov chain simulation from the statistical physics literature and the latent variable formula-

tion in maximum likelihood computation (i.e., EM algorithm) came together to spark the widespread application of MCMC methods in Bayesian computation.

## 2. SOME PRE-HISTORY

### 2.1 Markov Chain Monte Carlo

The origin of MCMC can be traced to the early 1950s when physicists were faced with the need to numerically study the properties of many particle systems. The state of the system is represented by a vector $x = (x_1, x_2, \ldots, x_n)$, where $x_i$ is the coordinate of the ith particle in the system and the goal is to study properties such as pressure and kinetic energy, which can be obtained from computation of the averaged values of suitably defined functions of the state vector. The averaging is weighted with respect to the canonical weight $\exp(-E(x)/kT)$, where the constants $k$ and $T$ denote the Boltzmann constant and the temperature, respectively. The physics of the system is encoded in the form of the energy function. For example, in a simple liquid model one has the energy $E(x) = (1/2) \sum \sum_{i \neq j} V(|x_i - x_j|)$, where $V(\cdot)$ is a potential function giving the dependence of pair-wise interaction energy on the distance between two particles. Metropolis et al. (1953) introduce the first Markov chain Monte Carlo method in this context by making sequential moves of the state vector by changing one particle at a time. In each move, a random change of a particle is proposed, say, by changing to a position chosen within a fixed distance from its current position, and the proposed change is either accepted or rejected according to a randomized decision that depends on how much the energy of the system is changed by such a move. Metropolis et al. justified the method via the concepts of ergodicity and detailed balance as in kinetic theory. Although they did not explicitly mention "Markov chain," it is easy to translate their formulation to the terminology of modern Markov chain theory. In subsequent development, this method was applied to a variety of physical systems such as magnetic spins, polymers, molecular fluids and various condense matter systems (reviewed in Binder, 1978), but all these applications share the characteristics that $n$ is large[1] and the $n$ components are homogeneous in the sense that each

takes value in the same space (say, 6-dimensional phase space, or up/down spin space, etc.) and interacts in identical manner with other components according to the same physical law as specified by the energy function. These characteristics made it difficult to recognize how the method can be of use in a typical Bayesian statistical inference problem where the form of the posterior distribution is very different from the Boltzmann distributions arising from physics. For this reason, although the probability and statistical community was aware of MCMC very early on (Hammersley and Handscomb, 1964[2]) and had in fact made key contributions to its theoretical development (Hastings, 1970), the method was not applied[3] to Bayesian inference until the 1980s.

---

[1] In the words of Geman and Geman (1984), "The Metropolis algorithm and others like it were invented to study the equilibrium properties, especially ensemble averages, time-evolution and low-temperature behavior, of very large systems of essentially identical, interacting components, such as molecules in a gas or atoms in binary alloys."

[2] Interestingly, Hammersley and Handscomb (1964) present the discrete version of the Metropolis algorithm in the chapter entitled "Problems in statistical mechanics," rather than in the chapter on "Principles of the Monte Carlo method" which covers topics such as crude Monte Carlo, stratified sampling, importance sampling, control variates, antithetic variates and regression methods. Similarly, in another popular textbook, Rubinstein (1981) presents the discrete version of the Metropolis algorithm toward the end of his book, embedded in an algorithm for 'global' optimization, rather than earlier in, for example, Chapter 3, "Random variate generation" or in Chapter 4, "Monte Carlo integration and variance reduction techniques." Also, Ripley (1987) presents the discrete version of the Metropolis algorithm in the section "Metropolis' method and random fields" (but not in the previous chapter on simulating random variables or in the following chapter on Monte Carlo integration and importance sampling).

[3] Several authors have argued that this delay can be attributed to "lack of appropriate computing power" (quoting Robert and Casella, 2010). Kass (1997) in his *JASA* review of the book *Markov Chain Monte Carlo in Practice* remarks on page 1645: "I believe that MCMC became important in statistics when it did for the simple reason that in the early 1990s large numbers of researchers could implement it on their desktops for interesting, nontrivial problems... I would suggest that the timing of their growth in popularity is explained primarily by computing technology." [This argument is quoted almost verbatim by Hitchcock (2003)—see Gubernatis (2003) who presents several "first-hand" accounts of the history of the Metropolis algorithm, as well as discusses why the algorithm received scant use in the 10–15 years following its development.]

This response does not explain why in 1987, when research universities had long since transitioned away from the university mainframe (which billed for every second of CPU usage, every page of memory, etc.) to departmental workstations (typically VAX 780's or 750's, which were essentially free to

## 2.2 Latent Variable Methods in Likelihood Inference: The EM Algorithm

During the 1960s and 1970s, statisticians developed an approach to maximum likelihood computation that is quite effective in many popular statistical models. The approach was based on the introduction of latent variables into the problem so as to make it feasible to compute the MLE if the latent variable value were known. Equivalently, if one regards the latent variable as "missing data", then this approach relies on the simplicity of inference based on the "complete data" to design an iterative algorithm to compute the maximum likelihood estimate and the associated standard errors. This development culminated in the publication of the extremely influential paper Dempster, Laird and Rubin (1977). A review of earlier research treating specific examples was presented in that paper, as well as the associated discussion. The high impact of Dempster, Laird and Rubin stems from its compelling demonstration that a wide variety of seemingly unrelated problems in standard statistical in-

---

faculty), the key computational Bayesian methods advocated at the time were based on quadrature or Laplace approximations or importance sampling (see the Appendix). Of course, one could not carry a VAX 750 around in a briefcase and it was not as powerful as current PCs, but one could log on from home or from the office and submit a job to run overnight (or over a weekend) without having to deal with a charge-back system. Clearly, the state of computing in the early 1980s did not detract from the work of Geman and Geman (1984).

It may not have been a lack of computing power *per se* that contributed to this developmental delay, but possibly a concern regarding optimal algorithms. In this regard, Geman (1988a) notes: "... we view Monte Carlo optimization techniques as *research tools*. They are poor substitutes for the efficient dedicated algorithms that should be developed when facing applications involving a flow of data and a need for speedy analysis." Efron (1979) mentions that Tukey wanted to call the *bootstrap* the *shotgun*, because it "can blow off the head of any problem if the statistician can stand the resulting mess." The same can be said for Metropolis–Hastings algorithms, as they can solve any simulation problem, but as noted by Geman (1988a), there is a strong sense that one could find a more efficient approach to implementing the simulation *in the specific case.*

Robert and Casella (2010) also cite "lack of background on Markov chains." We question this as well, for by the mid-1970s it was standard practice to require a stochastic process course based on such texts as Karlin and Taylor (1975) as part of the degree requirements for a Ph.D. in Statistics. The treatment in Karlin and Taylor (1975), for example, provides a more than sufficient background to understand the related material in Hammersley and Handscomb (1964).

---

ference, including multinomial sampling, normal linear models with missing values, grouping and truncation, mixture problems and hierarchical models, can all be encompassed within this latent variable framework and thus become computationally feasible using the same algorithm (called the EM algorithm by Dempster, Laird and Rubin) for MLE inference.

Because of its influence in later MCMC methods on the same set of problems, we briefly review a simplified formulation of the EM algorithm: Let $y$ be the observed data vector, $p_\theta(y)$ be the density of $y$, and we are interested in the inference regarding $\theta$. Two conditions are assumed for the application of the EM. First, it is assumed that although the likelihood $L(\theta|y) = p_\theta(y)$ may be hard to work with, one can introduce a latent (i.e., unobserved) variable $z$ so that the likelihood $L(\theta|y,z) = p_\theta(y,z)$ based on the value of $y$ and $z$ becomes easy to optimize as a function of $\theta$. In fact, for simplicity, we assume that $p_\theta(y,z)$ is an exponential family distribution. The second condition is that for any fixed parameter value $\theta$, it is possible to compute the expectation of the sufficient statistics of the exponential family, where the expectation is over $z$ under the assumption that $z$ is distributed according to its conditional distribution $p_\theta(z|y)$. We will see below that these conditions are closely related to the ones under which the most popular form of MCMC algorithm for Bayesian computation, namely, the Gibbs sampler, is applicable.

## 3. EMERGENCE OF MCMC IN BAYESIAN COMPUTATION IN THE 1980S

### 3.1 The State of Bayesian Computation in the Mid-1980s

The early 1980s was an active period in the development of Bayesian computational methods. In addition to the traditional approach that relied on the use of conjugate priors to obtain analytically tractable posterior distributions, significant progress was made in numerical approximations to the posterior distribution. We now briefly review the main approaches.

In many problems it is easy to evaluate the joint posterior density up to a constant of proportionality. The difficulty is to obtain posterior moments and marginal distributions of selected parameters of interest. Numerical integration methods were developed to obtain these quantities from the joint

posterior. In particular, Naylor and Smith (1982) and Smith et al. (1985) advocated the use of Gaussian quadrature which would be the correct choice in large sample situations when the posterior is approximately normal. Alternatively, Kloek and van Dijk (1978, 1980) proposed the use of importance sampling to carry out the integration, and applied the method systematically in the context of simultaneous equation models. Many novel variations were experimented with in both approaches, including the important idea of adaptation where a preliminary integration was used to guide the choice of parameters (grid points, importance function, etc.) in a second round of integration.

Another influential work was Tierney and Kadane (1986) which uses the technique of Laplace approximation to obtain accurate approximations for posterior moments and marginal densities; albeit in contrast to the other approaches, the accuracy of this approximation is determined by the sample size and not under the control of the Bayesian analyst.

These efforts demonstrated that accurate numerical approximation to marginal inference can be obtained in problems with moderate dimensional parameter space [e.g., Smith et al. (1985) report success on problems with up to 6 dimensions] and created a great deal of excitement in the prospect of computational Bayesian inference (see, e.g., Zellner, 1988). On the other hand, a review of the writings of leading Bayesian statisticians in this period reveal no awareness of the promise of the MCMC approach that would soon emerge as a dominate tool in Bayesian computation. In fact, among more dogmatic Bayesians,[4] the use of Monte Carlo was met with resistance and viewed as antithetical to Bayesian principles.

We are now ready to discuss the specific developments that sparked the emergence of MCMC methodology in statistics.

### 3.2 Formulation of the Gibbs Sampler

In 1984, Geman and Geman published a paper on the topic of Bayesian image analysis (Geman and Geman, 1984). Beyond its immediate and large impact in image analysis, this paper is significant for several results of more general interest, including a proof of the convergence of simulated annealing, and the introduction of the Gibbs sampler.

We briefly review how the Gibbs sampler emerged in this context. The authors began by drawing an analogy between images and statistical mechanics systems. Pixel gray levels and edge elements were regarded as random variables and an energy function based on local characteristics of the image was used to represent prior information on the image such as piece-wise smoothness. Because interaction energy terms involved only local neighbors, the conditional distribution of a variable given the remaining components of the image depends only on its local neighbors, and is therefore easy to sample from. Such a distribution, for the systems of image pixels, is similar to the canonical distribution in statistical mechanics studied by Boltzmann and Gibbs, and it is thus called a Gibbs distribution for the image.

Next, the authors analyzed the statistical problem of how to restore the image from an observed image which is a degradation of true image through the processes of local blurring and noise contamination. They showed that the posterior distribution of true image given the observed data is also a Gibbs distribution whose energy function still involves only local interactions. Geman and Geman proposed to generate images from this posterior distribution by iteratively sampling each image element from its conditional distribution given the rest of the image, which is easy to do because the distribution is still Gibbs. They call this iterative conditional sampling algorithm the Gibbs sampler. For the history of Bayesian computation, this was a pivotal step—although similar algorithms were already in use in the physics literature; to our knowledge, this work represented the first proposal to use MCMC to simulate from a posterior distribution. On the other hand, because the Gibbs model for images is so similar to the (highly specialized) statistical physics models, it was not apparent that this approach could be effective in traditional statistical models (see the Appendix).

### 3.3 Introduction of Latent Variables and Collective Moves

The use of iterative sampling for Bayesian inference in traditional statistical models was first demonstrated in Tanner and Wong (1987). The problems treated in this work, such as normal covariance estimation with missing data, latent class models, etc., were of the type familiar to mainstream statisticians of the time. A characteristic of many of these

---

[4] For a more detailed review of the prevailing approaches and views on Bayesian computation through the late 1980s, see the Appendix.

problems was that the likelihood is hard to compute (thus not amenable to MCMC directly). To perform Bayesian analysis on these models, the authors embedded them in the setting of the EM algorithm where a latent variable $z$ can be introduced to simplify the inference of the parameter $\theta$. They started from the equations

$$(3.1) \qquad p(\theta|y) = \int p(\theta|y,z)p(z|y)\,dz,$$

$$(3.2) \qquad p(z|y) = \int p(z|\theta,y)p(\theta|y)\,d\theta.$$

Recall that the conditions needed for the EM to work well are that $p_\theta(y,z)$ is simple to work with as a function of $\theta$, and $p_\theta(z|y)$ is easy to work with as a function of $z$. The first condition usually implies that the complete data posterior $p(\theta|y,z)$ is also easy to work with. Thus, (3.1) can be approximated as a mixture of $p(\theta|y,z)$ over a set of values (mixture values) for $z$ drawn from (3.2). Similarly, (3.2) is approximated as a mixture of $p(z|\theta,y)$ over mixture values for $\theta$ drawn from (3.1). This led the authors to propose an iterated sampling scheme to construct approximations to $p(\theta|y)$ and $p(z|y)$ simultaneously. In each step of the iteration, one draws a sample of values with replacement from the mixture values for $z$ (or $\theta$), and then conditional on each such $z$, draws $\theta$ (or $z$) from $p(\theta|y,z)$ [or $p(z|\theta,y)$].

This computation is almost identical to a version of the Gibbs sampler that iterates between the sampling of $p(\theta|y,z)$ and $p(z|\theta,y)$. In fact, if the sampling from the mixture values for $z$ (or $\theta$) was done without replacement rather than with replacement, as suggested by Morris (1987), then one would have exactly a population of independently run Gibbs samplers. The authors also noted[5] the connection to the equilibrium distribution of a Markov chain, but they did not employ it as the main mathematical framework in their analysis. In any case, a prominent aspect of its relevance lies in the explicit introduction of the latent variable $z$, which may or

may not be part of the data vector or the parameter vector of the original statistical model, to create an iterative sampling scheme for the Bayesian inference of the original parameter $\theta$. Tanner and Wong referred to this aspect of the design of the algorithm as "data augmentation." A judicious choice of latent variables can allow one to sample from the posterior $p(\theta|y)$ in cases where direct MCMC methods, including the Gibbs sampler, may not even be applicable because of difficulty in evaluating $p(\theta|y)$.

As a discussant of Tanner and Wong (1987), Morris (1987) makes several key observations of great relevance to MCMC Bayesian computing. In addition to suggesting a version of data augmentation that is the same as parallel Gibbs sampling, he emphasizes that (just as in the EM context) the augmentation is not limited to missing data, but can be done with parameters as well: "and to emphasize that their 'missing data' concept can be used to include unknown parameters or latent data." As an illustration of the data augmentation algorithm, Morris (1987) presents what we would now call the Gibbs sampler for a three-stage hierarchical model with $k+1$ parameters. At the first level of his model, $y_i|\theta_i$ are distributed independently as $N(\theta_i, V_i)$, for $i = 1,\ldots,k$ ($V_i$ known). At the second stage, $\theta_i|A$ are i.i.d. $N(0, A), i = 1,\ldots,k$. At the final stage, $A$ is distributed as a completely known distribution. Morris then says, "Let initial values $A^{(1)},\ldots,A^{(m)}$ be given. The posterior distribution of $\theta$ given $(y, A)$ is normally distributed and the P step samples $\theta_i^{(j)} \sim N((1-B_i^{(j)})y_i, V_i(1-B_i^{(j)}))$ for $i = 1,\ldots,k, j = 1,\ldots, m$, independently, with $B_i^{(j)} = V_i/(V_i + A^{(j)})\ldots$" For the $A$ parameter he writes, "The I step (1.5), therefore, samples new values $A^{(1)},\ldots,A^{(m)}$ according to $A^{(j)} \sim (\lambda + \|\theta^{(j)}\|^2)/\chi_{k+q}^2$ for $j = 1,\ldots,m$, with $\chi_{k+q}^2$ sampled independently for each $j$, $\|\theta\|^2$ denoting the sum of squares."

Although no new theory beyond MCMC is needed for the analysis of a sampling algorithm designed to include latent variables, the nature of the resulting process may be drastically different from the traditional MCMC processes, even in cases when the posterior $p(\theta|y)$ is computable and therefore amenable to direct MCMC analysis. Consider, for example, a linear model $y = x\beta + \varepsilon$, where the errors are independently distributed according to Student's $t$ with a fixed degree of freedom. In this case the joint posterior density for $\beta$ is computable and one can apply the standard Metropolis sampler to sample from it,

---

[5] Quoting Tanner and Wong (1987): "To see this, consider the extreme case in which $m = 1$, so that iteration $i$ produces only one value $\theta(i)$. In this case, $\theta(i)\,(i = 1, 2,\ldots)$ forms a Markov process with transition function equal to $K(\theta, \phi)$, as defined in (2.3). Under the regularity conditions of Section 6, this is an ergodic Markov process with an equilibrium distribution satisfying the fixed point equation given in (2.3)." Equation (2.3) presents the Markov transition function $K(\theta, \phi)$ (see also their Markov transition operator discussion in Remark 4) via the expression $g(\theta) = \int K(\theta, \phi)g(\phi)\,d\phi$, where $K(\theta, \phi) = \int p(\theta|z, y)p(z|\phi, y)\,dz$.

by iteratively proposing to change the $\beta$ vector, one component at a time. However, the sampler may be easily trapped at local maxima. What is worse, the moves may be exceedingly localized (therefore slow) if there is serious collinearity. On the other hand, by regarding the error as a gamma-mixture of normal variables, one can condition on the gamma variables and generate the whole vector $\beta$ (i.e., a collective move), which allows large moves even in the presence of collinearity.[6]

Interestingly, at about the same time, Swendsen and Wang (1987) also introduced the use of latent variables (called auxiliary variables) in the setting of a statistical mechanics system. This work deals with the Potts model of spins on a lattice. The authors introduced bond variables between spins and then alternated between the sampling of the two types of variables, spins and bonds. By conditioning on the bonds, they were able to make more global changes of the spin configuration by simultaneously updating a whole cluster of spins that are connected by active bonds (i.e., a collective move). In this way they were able to dramatically reduce the correlation time of the resulting Markov process for simulating a two-dimensional Ising model. Justifiably, this work is widely regarded as a breakthrough in dynamic Monte Carlo methods in statistical physics.

### 3.4 A Synthesis

Above we described how MCMC Bayesian computation arose in the 1980s from two independent sources, the statistical physics heritage as represented by Geman and Geman (1984), and the EM heritage as represented by Tanner and Wong (1987). A synthesis of these two traditions occurred in the important work of Gelfand and Smith (1990). Like the former, they employed the Gibbs sampling version of MCMC. Like the latter, they focused on traditional statistical models and relied on the use of latent variables to create iterative sampling schemes. Their paper[7] provided many examples to illustrate

the ease of use[8] and effectiveness of iterative sampling, and clarified the relation between the data augmentation algorithm and the Gibbs sampler.

The framing of data augmentation as MCMC also raised some new and interesting theoretical issues in the analysis of the MCMC output. For example, it follows from (3.2) that in data augmentation the estimate of an expectation $E(g(\theta)|y)$ is given by $\frac{1}{n}\sum_{i=1}^{n} E(g(\theta)|y, z_i)$, where the $z_i$'s are the currently sampled values for the latent variable $z$. Gelfand and Smith refer to the use of this estimate, instead of the usual estimate $\frac{1}{n}\sum_{i=1}^{n} g(\theta_i)$, as Rao–Blackwellization. They reasoned that if the $z_i$'s are independently drawn, as in a final iteration of the data augmentation algorithm, then clearly Rao–Blackwellization will reduce estimation error. They did not analyze the situation when the samples are dependent, as when the samples are generated from the Gibbs sampling process. The superiority of the Rao–Blackwellized estimator in the two-component Gibbs sampler was later established in Liu, Wong and Kong (1994)—see also Geyer (1995).

After the publication of Gelfand and Smith's influential paper, many mainstream statisticians began to adapt the use of MCMC in their own research, and the results in these early applications quickly established MCMC as a dominant methodology in Bayesian computation. However, it should be noted that in any given problem there could be a great many ways to formulate a MCMC sampler. In simulating an Ising model, for example, one can try to flip each spin conditional on the rest, or flip a whole set of spins connected by (artificially introduced) bonds that are sampled alternatively with

---

[6]For simplicity, the errors are unscaled in our example. If the errors have a scaling parameter, then it will have correlation with the gamma variables and this will slow down the convergence. More advanced augmentation schemes can be created to break this correlation and accelerate convergence; see Liu, Rubin and Wu (1998), Liu and Wu (1999) and van Dyk and Meng (2001).

[7]Also of note are the papers of Li (1988) who used a multi-component Gibbs sampler to perform multiple imputation from the posterior, Gelman and King (1990) who employed

MCMC to analyze hierarchical models of voting data across districts and Spiegelhalter and Lauritzen (1990) who treated graphical models. The latter appeared to have inspired the popular MCMC software BUGS (Lunn et al., 2009).

[8]Smith (1991) notes that regarding the Gibbs sampler: "Substantial iterative computation is then required, but the need for sophisticated numerical understanding on the part of the statistical analyst is obviated." Gelfand et al. (1990) also comment that sampling methods have a "hidden" efficiency: "In addition, in the case of most of the more sophisticated techniques, substantial fresh effort is required (including, in some cases, beginning the analysis anew) if the focus of inferential interest changes...." Tanner and Wong (1987, pages 533 and 549), in their normal variance–covariance matrix example, make the same point when they note the ease in which one may examine the posterior distribution of any function of the variance–covariance matrix, such as the smallest eigenvalue (see their Figure 1 in the Rejoinder).

the spins. The effectiveness of the Swendsen and Wang (1987) algorithm in the Ising model does not simply stem from the fact that it is a Gibbs sampler, but rather depends critically on the clever design of the specific form of the sampler. Likewise, a large part of the success of MCMC in the early 1990s was based on versions of Gibbs samplers that were designed to exploit the special structure of statistical problems in the style of the EM and data augmentation algorithms. Thus, the emergence of MCMC in mainstream Bayesian inference has depended as much on the introduction of the mathematically elegant MCMC formalism as on the realization that the structure of many common statistical models can be fruitfully exploited to design versions of the algorithm that are feasible and effective for these models.

The appearance of Gelfand and Smith (1990) marked the end of the emergence of the MCMC approach to the study of posterior distributions, and the beginning of an exciting period, lasting to this day, of the application of this approach to a vast array of problems, including inference in non-parametric problems. Advanced techniques have also been developed in this framework to accelerate convergence. Statisticians, no longer laggards in MCMC methodology, now rival physicists in the advancement of MCMC methodology.[9] It is our hope that this paper will serve as a useful historical context to understand current developments.

## APPENDIX

In this appendix we present three key resources that define the state-of-the-art Bayesian computing as of the mid to late 1980s.

### A.1 Smith et al.(1985)

A key reference that catalogs the tools in the Bayesian's armamentarium as of 1985 is the paper by Adrian F. M. Smith and colleagues entitled "The implementation of the Bayesian paradigm" (see Smith et al., 1985). After providing an overview of the goals of Bayesian computation, the authors critically review a number of implementation strategies: (1) exact analytic implementation based on conjugate priors; (2) large sample normal approximation; (3) alternative asymptotic approximations including discussion of a preprint of Tierney and Kadane (1986)

outlining the Laplace approach; (4) Monte Carlo integration—specifically importance sampling; (5) quadrature methodology—specifically the approach of Naylor and Smith (1982) based on the Gauss–Hermite product rules; and (6) successive transformations of the parameters of the model to achieve normality. The remainder of Smith et al. (1985) reviews in considerable detail the Naylor and Smith (1982) strategy based on Gaussian quadrature, as well as presents several interesting examples. The authors conclude with the following remark: "novel numerical integration techniques together with efficient graphical procedures are now making Bayesian analysis practical for a wide range of problems."

### A.2 The 1987 Special Issue of JRSS D (The Statistician)

In 1987, the *Journal of the Royal Statistical Society Series D (The Statistician)* published a special issue entitled "Practical Bayesian statistics," edited by Gopal K. Kanji with technical editors Adrian F. M. Smith and A. P. Dawid. Following up on the report of Smith et al. (1985), Smith et al. (1987) discuss an adaptive approach where Gauss–Hermite quadrature methods are combined with parameter transformations in an iterative manner, namely, successive transformations are determined by the estimated variance–covariance matrix of the previous iteration. Also discussed in this paper is an iterative importance sampling strategy, where the information in the previous iteration is used to improve the importance function for the transformed parameters. Smith et al. (1987) suggest using "quasi-random" sequences on the $k$-dimensional hypercube and reference the preprint of Shaw (1988a).

The paper by Smith et al. (1987) sets the computational theme of the issue, as many of the other papers make use of either Gauss–Hermite quadrature methodology or importance sampling. van Dijk, Hop and Louter (1987) present the details of an algorithm for the computation of posterior moments and densities based on importance sampling, specifically that the importance sampling function can be adapted based on the output of the previous iteration. In van Dijk, Hop and Louter (1987), the posterior mean and the posterior covariance matrix based on the output of the previous iteration is used to update the parameters of the multivariate Student $t$ importance function.

Stewart (1987) illustrates how (nonadaptive) importance sampling can be used in the context of

---

[9]For example, Geyer (1991) introduced parallel tempering ahead of related concepts in the physics literature—see Hukushima and Nemoto (1996).

hierarchical Bayesian models. Grieve (1987), Marriott (1987), Naylor (1987) and Shaw (1987) overcome the analytic intractability of the posterior using the adaptive Gauss–Hermite integration strategy of Naylor and Smith (1982). Van Der Merwe and Groenewald (1987) approximate the posterior distribution with a Pearson distribution, while Achcar, Bolfarine and Pericchi (1987) make use of the Laplace approximation discussed in Tierney and Kadane (1986). Spiegelhalter (1987), in his treatment of evidence propagation in expert systems, briefly mentions stochastic relaxation (Geman, 1988b) as a possible technique but does not use it in his treatment.[10] Finally, O'Hagan (1987) expressed strong objection to using Monte Carlo methods in Bayesian inference—even going as far as giving his paper the title "Monte Carlo is fundamentally unsound".

### A.3 The Third Valencia International Meeting: June 1–5, 1987

The Third Valencia International Meeting on Bayesian Statistics was held on June 1–5, 1987. According to the Preface of the Proceedings (see Bernardo et al., 1988), the scientific program consisted of 31 invited papers, each with discussion, and 33 refereed contributed papers: "The selection of topics, authors and discussants ensures that those Proceedings provide a definitive up-to-date overview of current concerns and activity in *Bayesian Statistics*, encompassing a wide range of theoretical and applied research."

As was seen in the paper by Smith et al. (1985), as well as the special issue of *The Statistician*, the key computational approaches vying for contention in this Proceedings are approximations based on Laplace's method, Gaussian quadrature numerical integration—possibly implemented in an iterative manner, and importance sampling—possibly embellished with an adaptive procedure to update the importance function. A survey of available software for Bayesian analysis was presented by Goel (1988) and this paper was discussed by van Dijk (1988). In the body of the discussion, van Dijk (1988) proposes a decision tree to determine which computational technique is most suitable for the problem at hand. He notes that if the posterior is "reasonably well behaved in the sense that it is unimodal, continuous, proper, not too skewed," then a Laplace approximation approach or possibly an importance sampling

approach using a Student $t$ importance function, as discussed in van Dijk, Hop and Louter (1987), should be used. If a transformation of the parameters results in a "more regular shape of the posterior," then he suggests Naylor and Smith (1982). If the posterior distribution is unimodal, but not much more is known, then van Dijk suggests importance sampling with adaptive importance functions. In this regard, van Dijk (1988) references the preprint of Geweke (1989) who also advocates for importance sampling and who remarks, "Integration by Monte Carlo is an attractive research tool because it makes numerical problems much more routine than do other numerical integration methods."

A second paper on Bayesian software was presented by Smith (1988). In the first sentence of the paper's abstract, he notes, "Recent developments in methods of numerical integration and approximation, in conjunction with hardware trends towards the widespread availability of single-user workstations which combine floating-point arithmetic power with sophisticated graphics facilities in an integrated interactive environment, would seem to have removed whatever excuses were hitherto historically available for the lack of any generally available form of Bayesian software." Smith argues that given advances in algorithm development, as well as the movement from large mainframes to workstations, the time is ripe for the development of Bayesian software. Smith points out on page 433 that the computational tools he has in mind are approximations based on asymptotic expansions, numerical integration methods based on quadrature and importance sampling methodology. In discussing the state of the field, Zellner (1988) notes, "... it is concluded that a Bayesian era in econometrics and statistics has emerged."

DuMouchel (1988), Albert (1988), Poirier (1988) and Sweeting (1988) employ or suggest the use of the Laplace approximation referencing Tierney and Kadane (1986) or related references. Morris (1988) proposes an approach based on the Pearson family which can be used to generalize the method of Laplace. Kass, Tierney and Kadane (1988) discuss the Laplace approximation in detail, extending the theory, providing examples and discussing implementation in the S computing environment.

Kim and Schervish (1988) make extensive use of the Gauss–Hermite approach as presented in Smith et al. (1985) and the discussants of this paper suggest the use of spherical integration rules (as implemented in *Bayes 4*). Shaw (1988b) discusses several

---

[10]See also Pearl (1987).

approaches to numerical integration, with emphasis on Gauss–Hermite and importance sampling with quasi-random sequences (see Shaw, 1987, 1988a). Grieve (1988) uses Gauss–Hermite in the analysis of $LD$50 experiments, Marriott (1988) uses these methods (referencing *Bayes 4*) in the context of ARMA time series models and Pole (1988) in the context of state-space models. Schnatter (1988) uses generalized Laguerre integration for forecasting $AR(p)$ time series models.

Rubin (1988) presents an overview[11] of his importance sampling based algorithm *SIR*, which he proposes as a general approach for posterior simulation, distinguishing it from the iterative MCMC approach in Tanner and Wong (1987), noting, "The SIR (Sampling/Importance Resampling) algorithm is an ubiquitously applicable noniterative algorithm for obtaining draws from an awkward distribution: $M$ draws from an initial approximation are made, and then $m < M$ draws are made from these with probability approximately proportional to their importance ratios."

A noted exception to the Laplace/numerical integration/importance sampling approach to Bayesian computing is the paper by Geman (1988a). In Section 2.5 on computing, Geman very clearly points out the basic idea behind MCMC methods, "Dynamics are simulated by producing a Markov chain, $X(1), X(2), \ldots$ with transition probabilities chosen so that the equilibrium distribution is the posterior (Gibbs) distribution (2.4). One way to do this is with the Metropolis algorithm (Metropolis et al., 1953). More convenient for image processing is a variation we call *stochastic relaxation*." He then proceeds to present what we now refer to as the full conditionals for the Gibbs sampler in the context of his problem. However, neither in the discussion nor in the response is there evidence to suggest that anyone at the time envisioned the MCMC methodology presented in Geman (1988a) having a broader impact to general parametric statistical problems. If anything, the methodology seems pigeon-holed as techniques for image analysis. Examining the Preface of the Proceedings, we find that the paper is summarized as "The important area of image-processing is reviewed by Geman." In fact, Geman appeared to have doubts about the practical utility of MCMC methods in his remark on page 169: "On the other hand, it is indeed difficult to find approaches that are as computationally expensive as ours. In this regard, we view Monte Carlo optimization techniques as *research tools*. They are poor substitutes for the efficient dedicated algorithms that should be developed when facing applications involving a flow of data and a need for speedy analysis." Added to this is his comment found on page 171 when referring to the algorithm in Besag (1986): "Furthermore, this *deterministic* algorithm is typically far more efficient than stochastic relaxation methods."

## ACKNOWLEDGMENTS

## REFERENCES

ACHCAR, J. A., BOLFARINE, H. and PERICCHI, L. R. (1987). Transformation of survival data to an extreme value distribution. *J. Roy. Statist. Soc. Ser. D Statistician* **36** 229–234.

ALBERT, J. (1988). Bayesian estimation of Poisson means using a hierarchical log-linear model. In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A. F. M. Smith, eds.) 519–531. Oxford Univ. Press, Oxford. MR1008064

BERNARDO, J. M., DEGROOT, M. H., LINDLEY, D. V. and SMITH, A. F. M., EDS. (1988). *Bayesian Statistics 3*. Oxford Univ. Press, Oxford. MR1008039

BESAG, J. (1986). On the statistical analysis of dirty pictures. *J. R. Stat. Soc. Ser. B Stat. Methodol.* **48** 259–302. MR0876840

BINDER, K. (1978). *Monte Carlo Methods in Statistical Physics*. Springer, New York. MR0555878

DEMPSTER, A. P., LAIRD, N. M. and RUBIN, D. B. (1977). Maximum likelihood from incomplete data via the EM algorithm (with discussion). *J. R. Stat. Soc. Ser. B Stat. Methodol.* **39** 1–38. MR0501537

DUMOUCHEL, W. (1988). A Bayesian model and a graphical elicitation procedure for multiple comparisons. In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A. F. M. Smith, eds.) 127–146. Oxford Univ. Press, Oxford. MR1008039

EFRON, B. (1979). Bootstrap methods: Another look at the Jackknife. *Ann. Statist.* **7** 1–26. MR0515681

GELFAND, A. E., HILLS, S. E., RACINE-POON, A. and SMITH, A. F. M. (1990). Illustration of Bayesian inference in normal data models using Gibbs sampling. *J. Amer. Statist. Assoc.* **85** 972–985.

GELFAND, A. E. and SMITH, A. F. M. (1990). Sampling-based approaches to calculating marginal densities. *J. Amer. Statist. Assoc.* **85** 398–409. MR1141740

---

[11]The presentation of this algorithm comprised the bulk of his discussion of Tanner and Wong (1987).

GELMAN, A. and KING, G. (1990). Estimating the electoral consequences of legislative redistricting. *J. Amer. Statist. Assoc.* **85** 274–282.

GEMAN, S. (1988a). Experiments in Bayesian image analysis. In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A. F. M. Smith, eds.) 159–171. Oxford Univ. Press, Oxford.

GEMAN, S. (1988b). Stochastic relaxation methods for image restoration and expert systems. In *Maximum Entropy and Bayesian Methods in Science and Engineering (Vol. 2)* (G. J. Erickson and C. R. Smith, eds.). Kluwer, New York. MR0970827

GEMAN, S. and GEMAN, D. (1984). Stochastic relaxation, Gibbs distributions and the Bayesian restoration of images. *IEEE Trans. Pattern Anal. Mach. Intell.* **6** 721–741.

GEWEKE, J. (1989). Bayesian inference in econometric models using Monte Carlo integration. *Econometrica* **57** 1317–1339. MR1035115

GEYER, C. J. (1991). Markov chain Monte Carlo maximum likelihood. In *Computing Science and Statistics: Proceedings of the 23rd Symposium on the Interface* (E. Keramidas, ed.) 156–163. Interface Foundation, Fairfax Station.

GEYER, C. J. (1995). Conditioning in Markov chain Monte Carlo. *J. Comput. Graph. Statist.* **4** 148–154. MR1341319

GOEL, P. K. (1988). Software for Bayesian analysis: Current status and additional need. In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A. F. M. Smith, eds.) 173–188. Oxford Univ. Press, Oxford.

GRIEVE, A. P. (1987). Applications of Bayesian software: Two examples. *J. Roy. Statist. Soc. Ser. D Statistician* **36** 283–288.

GRIEVE, A. P. (1988). A Bayesian approach to the analysis of LD50 experiments. In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A. F. M. Smith, eds.) 617–630. Oxford Univ. Press, Oxford.

GUBERNATIS, J. E., ed. (2003). *The Monte Carlo Method in the Physical Sciences: Celebrating the 50th Anniversary of the Metropolis Algorithm.* Amer. Inst. Phys., New York.

HAMMERSLEY, J. M. and HANDSCOMB, D. C. (1964). *Monte Carlo Methods*, 2nd ed. Chapman and Hall, London. MR0223065

HASTINGS, W. K. (1970). Monte Carlo sampling methods using Markov chains and their applications. *Biometrika* **57** 97–109.

HITCHCOCK, D. B. (2003). A history of the Metropolis–Hastings algorithm. *Amer. Statist.* **57** 254–257. MR2037852

HUKUSHIMA, K. and NEMOTO, K. (1996). Exchange Monte Carlo method and application to spin glass simulations. *J. Phys. Soc. Japan* **65** 1604–1608.

KARLIN, S. and TAYLOR, H. M. (1975). *A First Course in Stochastic Processes*, 2nd ed. Academic Press, New York. MR0356197

KASS, R. E. (1997). Review of "Markov chain Monte Carlo in practice." *J. Amer. Statist. Assoc.* **92** 1645–1646.

KASS, R. E., TIERNEY, L. and KADANE, J. B. (1988). Asymptotics in Bayesian computation. In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A. F. M. Smith, eds.) 261–278. Oxford Univ. Press, Oxford. MR1008051

KIM, C. E. and SCHERVISH, M. J. (1988). Stochastic models of incarceration careers. In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A. F. M. Smith, eds.) 279–305. Oxford Univ. Press, Oxford. MR1008052

KLOEK, T. and VAN DIJK, H. K. (1978). Bayesian estimates of equation system parameters: An application of integration by Monte Carlo. *Econometrica* **46** 1–19.

KLOEK, T. and VAN DIJK, H. K. (1980). Further experience in Bayesian analysis using Monte Carlo integration. *J. Econometrics* **14** 307–328.

LI, K. H. (1988). Imputation using Markov chains. *J. Statist. Comput. Simul.* **30** 57–79. MR1005883

LIU, C., RUBIN, D. B. and WU, Y. N. (1998). Parameter expansion to accelerate EM: The PX-EM algorith. *Biometrika* **85** 755–770. MR1666758

LIU, J. S., WONG, W. H. and KONG, A. (1994). Covariance structure of the Gibbs sampler with applications to the comparisons of estimators and augmentation schemes. *Biometrika* **81** 27–40. MR1279653

LIU, J. S. and WU, Y. N. (1999). Parameter expansion scheme for data augmentation. *J. Amer. Statist. Assoc.* **94** 1264–1274. MR1731488

LUNN, D., SPIEGELHALTER, D. J., THOMAS, A. and BEST, N. (2009). The BUGS project: Evolution, critique and future directions. *Stat. Med.* **28** 3049–3067.

MARRIOTT, J. (1987). Bayesian numerical and graphical methods for Box–Jenkins time series. *J. Roy. Statist. Soc. Ser. D Statistician* **36** 265–268.

MARRIOTT, J. (1988). Reparametrization for Bayesian inference in ARMA time series. In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A. F. M. Smith, eds.) 701–704. Oxford Univ. Press, Oxford.

METROPOLIS, N., ROSENBLUTH, A. W., ROSENBLUTH, M. N., TELLER, A. H. and TELLER, E. (1953). Equation of state calculations by fast computing machines. *J. Chem. Phys.* **21** 1087–1092.

MORRIS, C. N. (1987). Comment on "The calculation of posterior distributions by data augmentation" by M. A. Tanner and W. H. Wong. *J. Amer. Statist. Assoc.* **82** 542–543. MR0898357

MORRIS, C. N. (1988). Approximating posterior distributions and posterior moments. In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A. F. M. Smith, eds.) 327–344. Oxford Univ. Press, Oxford. MR1008054

NAYLOR, J. C. (1987). Bayesian alternatives to *t*-tests. *J. Roy. Statist. Soc. Ser. D Statistician* **36** 241–246.

NAYLOR, J. C. and SMITH, A. F. M. (1982). Applications of a method for the efficient computation of posterior distributions. *J. Roy. Statist. Soc. Ser. C Appl. Statist.* **31** 214–225. MR0694917

O'HAGAN, A. (1987). Monte Carlo is fundamentally unsound. *J. Roy. Statist. Soc. Ser. D Statistician* **36** 247–249.

PEARL, J. (1987). Evidential reasoning using stochastic simulation of causal models. *Artif. Intell.* **32** 245–257. MR0885357

POIRIER, D. J. (1988). Bayesian diagnostic testing in the general linear normal regression model. In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A.

F. M. Smith, eds.) 725–732. Oxford Univ. Press, Oxford. MR1008083

POLE, A. (1988). Transfer response models: a numerical approach. In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A. F. M. Smith, eds.) 733–745. Oxford Univ. Press, Oxford.

RIPLEY, B. D. (1987). *Stochastic Simulation*. Wiley, New York. MR0875224

ROBERT, C. and CASELLA, G. (2010). A short history of Markov chain Monte Carlo—subjective recollections from incomplete data. In *Handbook on Markov Chain Monte Carlo*. Chapman and Hall/CRC Press, Boca Raton, FL.

RUBIN, D. B. (1988). Using the SIR algorithm to simulate posterior distributions. In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A. F. M. Smith, eds.) 395–402. Oxford Univ. Press, Oxford.

RUBINSTEIN, R. Y. (1981). *Simulation and the Monte Carlo Method*, 1st ed. Wiley, New York. MR0624270

SCHNATTER, S. (1988). Bayesian forecasting of time series by Gaussian sum approximation. In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A. F. M. Smith, eds.) 757–764. Oxford Univ. Press, Oxford. MR1008084

SHAW, J. E. H. (1987). Numerical Bayesian analysis of some flexible regression models. *J. Roy. Statist. Soc. Ser. D Statistician* **36** 147–153.

SHAW, J. E. H. (1988a). A quasirandom approach to integration in Bayesian statistics. *Ann. Statist.* **16** 895–914. MR0947584

SHAW, J. E. H. (1988b). Aspects of numerical integration and summarisation. In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A. F. M. Smith, eds.) 411–428. Oxford Univ. Press, Oxford. MR1008059

SMITH, A. F. M. (1988). What should be Bayesian about Bayesian software? In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A. F. M. Smith, eds.) 429–435. Oxford Univ. Press, Oxford.

SMITH, A. F. M. (1991). Bayesian computational methods. *Philos. Trans. Roy. Soc. Lond. Ser. A* **337** 369–386. MR1143728

SMITH, A. F. M., SKENE, A. M., SHAW, J. E. H. and NAYLOR, J. C. (1987). Progress with numerical and graphical methods for practical Bayesian statistics. *J. Roy. Statist. Soc. Ser. D Statistician* **36** 75–82.

SMITH, A. F. M., SKENE, A. M., SHAW, J. E. H., NAYLOR, J. C. and DRANSFIELD, M. (1985). The implementation of the Bayesian paradigm. *Commun. Stat. Theory Methods* **14** 1079–1102. MR0797634

SPIEGELHALTER, D. J. (1987). Coherent evidence propagation in expert systems. *J. Roy. Statist. Soc. Ser. D Statistician* **36** 201–210.

SPIEGELHALTER, D. J. and LAURITZEN, S. L. (1990). Sequential updating of conditional probabilities on directed graphical structures. *Networks* **20** 579–605. MR1064738

STEWART, L. (1987). Hierarchical Bayesian analysis using Monte Carlo integration: Computing posterior distributions when there are many possible models. *J. Roy. Statist. Soc. Ser. D Statistician* **36** 211–219.

SWEETING, T. J. (1988). Approximate posterior distributions in censored regression models. In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A. F. M. Smith, eds.) 791–799. Oxford Univ. Press, Oxford. MR1008087

SWENDSEN, R. H. and WANG, J. S. (1987). Nonuniversal critical dynamics in Monte Carlo simulations. *Phys. Rev. Lett.* **58** 86–88.

TANNER, M. A. and WONG, W. H. (1987). The calculation of posterior distributions by data augmentation (with discussion). *J. Amer. Statist. Assoc.* **82** 528–550. MR0898357

TIERNEY, L. and KADANE, J. B. (1986). Accurate approximations for posterior moments and marginal densities. *J. Amer. Statist. Assoc.* **81** 82–86. MR0830567

VAN DER MERWE, A. J. and GROENEWALD, P. C. N. (1987). Bayes and empirical Bayes confidence intervals in applied research. *J. Roy. Statist. Soc. Ser. D Statistician* **36** 171–179.

VAN DIJK, H. K. (1988). Discussion of Goel. In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A. F. M. Smith, eds.) 187–188. Oxford Univ. Press, Oxford.

VAN DIJK, H. K., HOP, J. P. and LOUTER, A. S. (1987). An algorithm for the computation of posterior moments and densities using simple importance sampling. *J. Roy. Statist. Soc. Ser. D Statistician* **36** 83–90.

VAN DYK, D. A. and MENG, X. L. (2001). The art of data augmentation. *J. Comput. Graph. Statist.* **10** 1–50. MR1936358

ZELLNER, A. (1988). A Bayesian era. In *Bayesian Statistics 3* (J. M. Bernardo, M. H. Degroot, D. V. Lindley and A. F. M. Smith, eds.) 509–516. Oxford Univ. Press, Oxford. MR1008063