

# Introduction to Markov Chains for MCMC

*Jonathan Navarrete*

## Introduction

A (discrete) **Markov chain** is a random (stochastic) process  $X_t : t \in T$  that can take  $k$ -states over time  $t = 1, 2, \dots$ . The process starts at one of these states ( $i$ ) and proceeds forward to another state ( $j$ ) with some probability  $p_{i,j}$ , the probability of transitioning from state  $i$  to state  $j$ . The probabilities  $p_{i,j}$  are called *transition probabilities*. These transition probabilities can be structured into a matrix like

$$P_{m,n} = \begin{pmatrix} p_{1,1} & p_{1,2} & \cdots & p_{1,n} \\ p_{2,1} & p_{2,2} & \cdots & p_{2,n} \\ \vdots & \vdots & \ddots & \vdots \\ p_{m,1} & p_{m,2} & \cdots & p_{m,n} \end{pmatrix}$$

## Two-State Markov Chain

Consider an ordered stochastic process with  $k = 2$  states (or outcomes),  $[0, 1]$ . This chain takes  $X_1, X_2, \dots$  steps where at each step it transitions between the two states with some probability  $p_{i,j}$ . When the chain is at state  $i = 0$ , the chain can transition to state  $j = 1$  with some probability,  $p_{0,1} = \alpha$ , or remain at  $j = 0$  with probability  $p_{0,0} = 1 - \alpha$ .

Likewise, if the chain is currently at state  $i = 1$ , the chain can transition to  $j = 0$  with probability  $p_{1,0} = \beta$  or remain at  $j = 1$  with probability  $p_{1,1} = 1 - \beta$ .

In a Markov chain, the random variables  $X_t$  depend on the previous it state  $X_{t-1}$ . Thus, the random variables are not independent; autocorrelation exists. The random variables, albeit dependent, are only dependent on the previous iteration such that the transition probabilities only depend on the most recent state

$$P(X_t = x_t | X_{t-1} = x_{t-1}, X_{t-2} = x_{t-2}, \dots) = P(X_t = x_t | X_{t-1} = x_{t-1})$$

This conditional dependence on one-previous-step is known as the **Markov Property**.

From the conditional probabilities, we can form a **transition matrix**

$$\mathbf{P}_{2 \times 2} = \begin{matrix} & \begin{matrix} 0 & 1 \end{matrix} \\ \begin{matrix} 0 \\ 1 \end{matrix} & \begin{pmatrix} p_{0,0} & p_{0,1} \\ p_{1,0} & p_{1,1} \end{pmatrix} \end{matrix} = \begin{matrix} & \begin{matrix} 0 & 1 \end{matrix} \\ \begin{matrix} 0 \\ 1 \end{matrix} & \begin{pmatrix} p_{0,0} & p_{0,1} \\ p_{1,0} & p_{1,1} \end{pmatrix} \end{matrix}$$

where the  $i^{th}$  row gives the conditional distribution of  $X_t | X_{t-1}$ , and each row's probabilities sum up to 1.

## Example: 2-State Markov Chain

If from state  $i = 0$ , we have a probability of 0.4 to transition to  $j = 1$ ,  $P(X_t = 1|X_{t-1} = 1) = 0.4$ . And from state  $i = 1$  we have a probability of 0.55 of transitioning to state  $j = 0$ ,  $P(X_t = 1|X_{t-1} = 0) = 0.55$ . Then for  $n = 1, 2, \dots$  iterations, then our matrix would be

$$\mathbf{P}_{2 \times 2} = \begin{pmatrix} p_{0,0} & p_{0,1} \\ p_{1,0} & p_{1,1} \end{pmatrix} = \begin{pmatrix} 0.6 & 0.4 \\ 0.55 & 0.45 \end{pmatrix}$$

Then as a simulation exercise, we would have

```
m = 2000;
n = 1:m;


x = numeric(m);

x[1] = 0

alpha = 0.4;
beta = 0.55
# Simulation
for (i in 2:m){
  if (x[i-1]==0){
    x[i] = rbinom(1, 1, alpha)
  } else {
    x[i] = rbinom(1, 1, 1 - beta)
  }
}

## First 35 transitions of the chain
plot(x[1:35], type="b",
     main = "First 35 transitions of the chain",
     xlab="Step", ylab="State")
```

### First 35 transitions of the chain



As can be seen in the first  $t = 1, \dots, 35$  transitions of the chain, the Markov chain does not always leave its current state. It transitions with a given transition probability.

## Another Example: 3-state Markov chain

For this following example we'll be utilizing the `markovchain` library available from **CRAN**. Revolution Analytics (now an extension of Microsoft) has a nice blog post and tutorial on the use of this package. See *Getting Started with Markov Chains*

Consider a basketball team currently participating in a tournament. As they proceed through their games, assume the team has three states: *win*, *lose*, *tie*. Say, at the start of the season, they win their first game,  $X_0 = \text{win}$ . From this point, they have probabilities of winning, losing or drawing their next game.

For our favorite team, say these probabilities are  $P(X_t = \text{win} | X_{t-1} = \text{win}) = 0.35$ ,  $P(X_t = \text{lose} | X_{t-1} = \text{win}) = 0.45$ ,  $P(X_t = \text{tie} | X_{t-1} = \text{win}) = 0.2$ .

However, if they were to lose their first game, the probabilities change to  $P(X_t = \text{win} | X_{t-1} = \text{lose}) = 0.60$ ,  $P(X_t = \text{lose} | X_{t-1} = \text{lose}) = 0.30$ ,  $P(X_t = \text{tie} | X_{t-1} = \text{lose}) = 0.1$ .

And if they were to tie their first game, then  $P(X_t = \text{win} | X_{t-1} = \text{tie}) = 0.55$ ,  $P(X_t = \text{lose} | X_{t-1} = \text{tie}) = 0.40$ ,  $P(X_t = \text{tie} | X_{t-1} = \text{tie}) = 0.05$ . Thus, these probabilities construct a transition matrix  $\mathbf{P}$ .

$$P_{3 \times 3} = \begin{pmatrix} p_{\text{win},\text{win}} & p_{\text{win},\text{lose}} & p_{\text{win},\text{tie}} \\ p_{\text{lose},\text{win}} & p_{\text{lose},\text{lose}} & p_{\text{lose},\text{tie}} \\ p_{\text{draw},\text{win}} & p_{\text{tie},\text{lose}} & p_{\text{tie},\text{tie}} \end{pmatrix} = \begin{pmatrix} 0.35 & 0.45 & 0.20 \\ 0.60 & 0.30 & 0.10 \\ 0.55 & 0.40 & 0.05 \end{pmatrix}$$

```
library(markovchain)
```

```
## Package:  markovchain
## Version:  0.6.9.3
## Date:     2017-05-08
## BugReport: http://github.com/spedygiorgio/markovchain/issues


## set up transition matrix
P <- matrix(data = c(0.35, 0.45, 0.20,
                    0.60, 0.30, 0.10,
                    0.55, 0.40, 0.05), byrow = TRUE,
            nrow = 3, ncol = 3,
            dimnames = list(c("win", "lose", "tie"), c("win", "lose", "tie")))

P <- new( "markovchain", transitionMatrix = P)

print(P)

##      win lose tie
## win  0.35 0.45 0.20
## lose 0.60 0.30 0.10
## tie  0.55 0.40 0.05

plot(P)
```



As in the 2-state example, our Markov chain's transition to  $X_t$  only depends on  $X_{t-1}$  and no further. So for the probability outcome of the current game, we would only need to look back at the most recent past game and not the entire season.

## Limiting Probabilities

The *Chapman-Kolmogorov* equations provide a method for computing 1, 2, ...,  $n$ -step transition probabilities. The following can be used to calculate the  $n$ -step probabilities

$$P_{ij}^{n+m} = \sum_{x=0}^{\infty} P_{ix}^n P_{xj}^m \text{ for all } n, m \geq 0, \text{ all } i, j$$

Let  $\mathbf{P}^{(n)}$  denote the  $n$ -ste transition matrix, then by the *Chapman-Kolmogorov* equations

$$\mathbf{P}^{(n+m)} = \mathbf{P}^{(n)} \cdot \mathbf{P}^{(m)}$$

Thus, we can use simple matrix multiplication to evaluate the probability matrix for step  $n$ .

```
P <- matrix(data = c(0.35, 0.45, 0.20,
                    0.60, 0.30, 0.10,
                    0.55, 0.40, 0.05), byrow = TRUE,
            nrow = 3, ncol = 3,
            dimnames = list(c("win", "lose", "tie"), c("win", "lose", "tie")))

print(P) ## initial probabilities
```

```
##      win lose tie
## win  0.35 0.45 0.20
## lose 0.60 0.30 0.10
## tie  0.55 0.40 0.05
```

```
P_2 = P %*% P
print(P_2)
```

```
##      win  lose  tie
## win  0.5025 0.3725 0.1250
## lose 0.4450 0.4000 0.1550
## tie  0.4600 0.3875 0.1525
```

```
P_4 = P_2 %*% P_2
print(P_4)
```

```
##      win      lose      tie
## win  0.4757688 0.3846188 0.1396125
## lose 0.4729125 0.3858250 0.1412625
## tie  0.4737375 0.3854438 0.1408188
```

```
P_8 = P_4 %*% P_4
print(P_8)
```

```
##      win      lose      tie
## win  0.4743866 0.3851979 0.1404155
## lose 0.4743798 0.3852007 0.1404195
## tie  0.4743818 0.3851999 0.1404183
```

```
P_16 = P_8 %*% P_8
print(P_16)
```

```
##      win      lose      tie
## win  0.4743833 0.3851992 0.1404175
## lose 0.4743833 0.3851992 0.1404175
```

```
## tie 0.4743833 0.3851992 0.1404175
```

```
P_20 = P_16 %*% P_4  
print(P_20)
```

```
##           win      lose      tie  
## win 0.4743833 0.3851992 0.1404175  
## lose 0.4743833 0.3851992 0.1404175  
## tie 0.4743833 0.3851992 0.1404175
```

By step 20, it appears that we see a convergence in the  $n$ -step transition matrix. We've arrived at a limiting distribution for our 3-state Markov chain.

Not all Markov chains have limiting distributions, especially if there is not free movement across all states.

For a more thorough review of Markov chains, see *Chapter 11: Markov Chains*.



## Ergodic Properties

For MCMC computations, only Markov chains with specific properties can be used. Markov chains are constructed from a *transition kernel*  $K$ , a conditional probability density such that  $X_{t+1} \sim K(X_t, X_{t+1})$ . A typical example is provided by the random walk process, formally defined as follows. As we continue with notes and definitions, note that this will be through the perspective of a discrete-time stochastic processes,  $(X_t)_{t \in \mathbb{N}}$

### Random Walk

A sequence of random variables  $(X_t)$  is a *random walk* if it satisfies

$$X_{t+1} = X_t + \epsilon_t,$$

where  $\epsilon_t$  is generated independently of  $X_t, X_{t-1}, \dots$ . If the distribution of the  $\epsilon_t$  is symmetric about zero, the sequence is called a symmetric random walk.

There are many examples of random walks, and random walks play a key role in many MCMC algorithms, particularly those based on the Metropolis-Hastings algorithm.

The chains encountered in MCMC settings enjoy a very strong stability property, a *stationary probability distribution*. A stationary probability distribution exists if the kernel  $K$  allows for free movement all over the state space for all  $X_1, X_2, \dots, X_t$  transitions. This freedom to move all around the state space is called *irreducibility*, and is essential for MCMC algorithms. Irreducibility states that for  $n \in \mathbb{N}$  such that  $P(X_n \in A | X_0) > 0$  for every  $A$  such that  $\pi(A) > 0$ .

The irreducibility property also ensures that most of the chains involved in MCMC algorithms are *recurrent*, that the number of times a chain visits an arbitrary set  $A$  in the state space is infinity; a chain is allowed to revisit any part of the state space always.

The *stationary distribution* is also a *limiting distribution* in the sense that the limiting distribution of  $X_{t+1}$  is  $f$  under the total variation norm, regardless of the starting value  $X_0 = x_0$ . Thus, what we can observe from our MCMC simulations is that, as a result of these convergence properties, is that the empirical average

$$\frac{1}{T} \sum_{n=1}^T h(X_t) \rightarrow E_f[h(X)]$$

convergence to the expectation  $E[h(X)]$  almost surely.

Another condition to be reviewed is the *detailed balance condition*, also known as *reversibility*.

$$f(x)K(y|x) = f(y)K(x|y)$$

When the chain is *reversible* (that is, when the transition kernel is symmetric), a Central Limit Theorem also holds for this average. In summary, the Law of Large Numbers that lies at the basis of previous *i.i.d.* Monte Carlo methods can also be applied in MCMC settings.

Reversibility is important because it has the effect of balancing movement through the entire state space. When a Markov chain is reversible,  $f(\cdot)$  is the unique, invariant, stationary distribution of that chain. Hence, if  $f(\cdot)$  is of interest, we need only find the reversible Markov chain for which  $f(\cdot)$  is the limiting distribution. This is why MCMC works!

These conditions together form what is often referred to the **Ergodic Theorem**. That said, we can now begin delving into MCMC algorithms such as the Metropolis-Hastings algorithm or the Gibbs sampler, which by the ergodic theorem are almost always theretically convergent.