

# Gibbs Samplers

Jonathan Navarrete

# Introduction

Now that we've familiarized ourselves with MCMC and the Metropolis-Hastings algorithms, we begin to analyze a now-common MCMC algorithm called Gibbs sampler. The Gibbs sampler is in fact a special case of the Metropolis-Hastings algorithm for high dimensional target distributions.

Recommended reading: [Explaining the Gibbs Sampler](#)

# Two-stage Gibbs sampler algorithm

The Gibbs sampler creates a Markov chain from a joint distribution in the following way. If two random variables  $X$  and  $Y$  have joint density  $f(x, y)$ , with corresponding conditional densities  $f_{Y|X}$  and  $f_{X|Y}$ , the two stage Gibbs sampler generates a Markov chain  $(X_t, y_i)$  by generating  $y_i$  from conditional density  $f_{Y|X}$  and then generating  $X_t$  from conditional density  $f_{X|Y}$ .

The **two-stage Gibbs sampler algorithm** as described Robert & Casella goes as follows

Take  $X_0 = x_0$

For  $t = 1, 2, \dots$ , generate

1.  $y_i \sim f_{Y|X}(\cdot | x_{t-1})$

2.  $X_t \sim f_{X|Y}(\cdot | y_t)$

# Bivariate Normal Example

We illustrate the implementation of the Gibbs sampler with a simple example. Consider a bivariate Normal distribution where

$$X, Y \sim N_2 \left( \begin{pmatrix} \mu_X \\ \mu_Y \end{pmatrix}, \begin{pmatrix} \sigma_X^2 & \sigma_{XY}^2 \\ \sigma_{YX}^2 & \sigma_Y^2 \end{pmatrix} \right)$$

# Bivariate Normal Example

The marginal distributions of  $X$  and  $Y$  are  $N(\mu_X, \sigma_X)$  and  $N(\mu_Y, \sigma_Y)$ . The conditional distributions of  $Y$  and  $X$  are

$$Y|X = x \sim N\left(\mu_Y + \frac{\rho\sigma_Y}{\sigma_X}(x - \mu_X), (1 - \rho^2)\sigma_Y^2\right)$$

and

$$X|Y = y \sim N\left(\mu_X + \frac{\rho\sigma_X}{\sigma_Y}(y - \mu_Y), (1 - \rho^2)\sigma_X^2\right)$$

$\rho$  is the correlation between  $X$  and  $Y$ , and  $(1 - \rho^2)\sigma_X^2$  is the variance.

[Multivariate Normal Distribution](#)

# Bivariate Normal Example

```
library(MASS, quietly = TRUE) ## for kernel density estimation
library(RColorBrewer, quietly = TRUE) ## some pretty colors
k <- 11
my.cols <- rev(brewer.pal(k, "RdYlBu"))

N = 100000
rho = 0.9
mu_x = 1
mu_y = 2
sd_x = 1.2
sd_y = 0.75

s1 = sqrt(1 - rho^2) * sd_x
s2 = sqrt(1 - rho^2) * sd_y

MVN = matrix(data = NA, nrow = N, ncol = 2,
              dimnames = list(NULL, c("X", "Y")))
```

# Bivariate Normal Example

```
MVN[1, ] = c(mu_x, mu_y)
Y = MVN[1, 2] ## get Y vals
for(i in 1:(N-1)){
  mx = mu_x + rho * (Y - mu_y) * sd_x/sd_y
  X = rnorm(n = 1, mx, s1)
  MVN[i+1, 1] = X
  my = mu_y + rho * (X - mu_x) * sd_y/sd_x
  Y = rnorm(n = 1, mean = my, sd = s2)
  MVN[i+1, 2] = Y
}
```

```
## means
colMeans(MVN)
```

```
##           X           Y
## 1.000312 1.999968
```

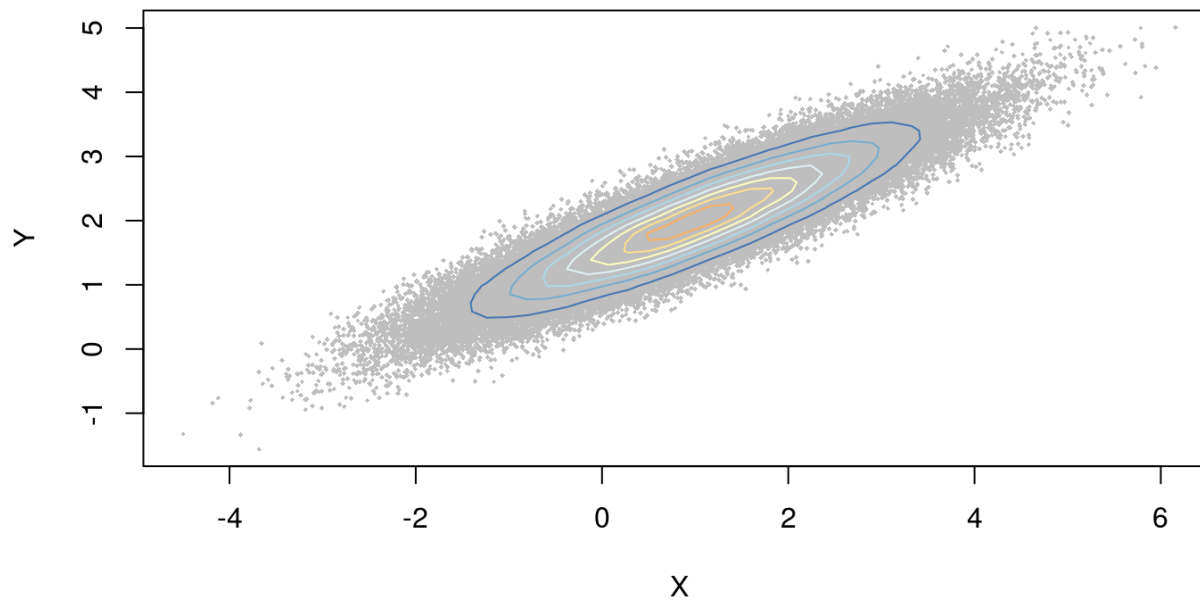
```
## correlation
cor(MVN)
```

```
##           X           Y
## X 1.0000000 0.9011088
## Y 0.9011088 1.0000000
```

# Bivariate Normal Example

```
z = kde2d(x = MVN[,1], y = MVN[,2], n = 50) ## kernel density estimation
plot(MVN, xlab="X", ylab="Y", pch = 18, col = 8, cex=.4, main = "MVN Samples")
contour(z, drawlabels=FALSE, nlevels=k, col=my.cols, add=TRUE) ## add contours to plot
```

**MVN Samples**





# Beta-Binomial revisited

In the introduction to these notes, we saw a Bayesian example of the Beta-Binomial distribution. From Casella's paper , we revisit this example.

$$X|\theta \sim \text{Bin}(n, \theta), \text{ and } \theta \sim \text{Beta}(a, b)$$

have joint density

$$f(x, \theta) = \binom{n}{x} \frac{\Gamma(a+b)}{\Gamma(a)\Gamma(b)} \theta^{x+a-1} (1-\theta)^{n-x+b-1}$$

is the a  $\text{Beta}(x+a, n-x+b)$  distribution.

# Beta-Binomial revisited

Suppose we are interested in calculating some characteristics of the marginal distributions of  $X|\theta$  and  $\theta|a, b, x, n$ .

$$\begin{aligned}f(x|\theta) &\text{ is } \textit{Bin}(n, \theta) \\f(\theta|x) &\text{ is } \textit{Beta}(x + a, n - x + b)\end{aligned}$$

Therefore, we follow an iterative algorithm of

$$\begin{aligned}X_i &\sim f(x|\theta) \\ \theta_{i+1} &\sim f(\theta|X_i = x_i)\end{aligned}$$

# Beta-Binomial revisited

```

n = 16
a = 2
b = 4
X = numeric(N)
thetas = numeric(N)

X[1] = runif(1) ## initial values
thetas[1] = runif(1) ## theta values

for(i in 1:N){
  X[i+1] = rbinom(n = 1, size = n, prob = thetas[i])
  thetas[i+1] = rbeta(1, a + X[i+1], n - X[i+1]+b)
}

quantile(X)

```

```

##      0%   25%   50%   75%  100%
##      0     3     5     8    16

```

```
quantile(thetas)
```

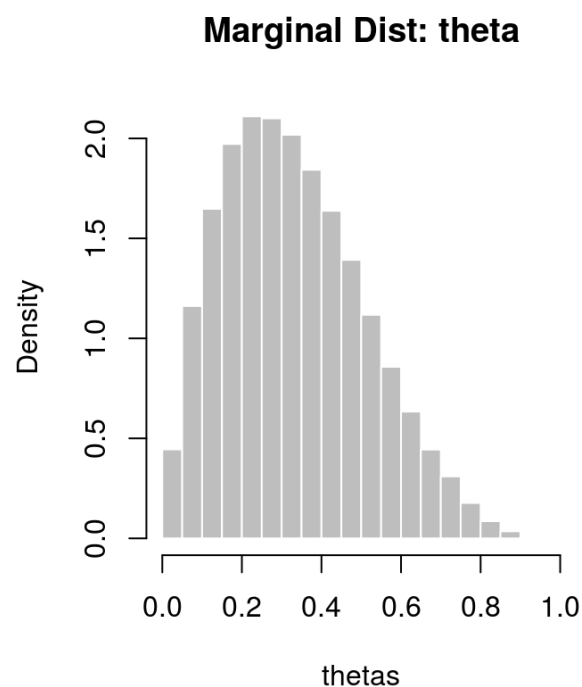
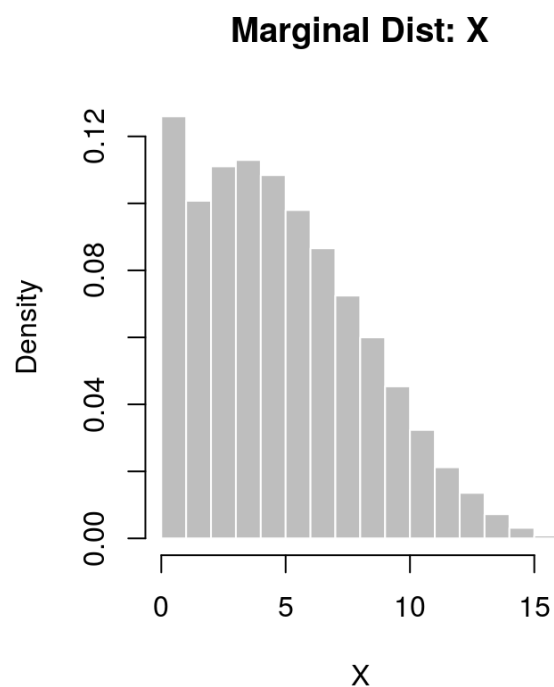
```

##              0%              25%              50%              75%              100%
## 0.0005819406 0.1944846157 0.3136805137 0.4520986182 0.9659229861

```

# Beta-Binomial revisited

```
par(mfrow = c(1,2))  
hist(X, main = "Marginal Dist: X", probability = TRUE, border = "white", col = "gray")  
hist(thetas, main = paste("Marginal Dist:", expression(theta)), probability = TRUE, border = "white", col = "gray")
```



# The multistage Gibbs sampler

There is a natural extension from the two-stage Gibbs sampler to the general multistage Gibbs sampler. Suppose that, for some  $p > 1$ , the random variable  $\mathbf{X} \in X$  can be written as  $\mathbf{X} = (X_1, \dots, X_p)^T$  where the  $X_i$ 's are either unidimensional or multidimensional components. Suppose that we can simulate from the corresponding conditional densities,  $f_1, f_2, \dots, f_p$  that is, we can simulate

$$X_i | x_1, \dots, x_{i-1}, x_{i+1}, \dots, x_p \sim f(x_i | x_1, \dots, x_{i-1}, x_{i+1}, \dots, x_p)$$

for  $i$  in  $1, 2, \dots, p$ . The associated Gibbs sampler is given as

# The multistage Gibbs sampler

At iteration  $t = 1, 2, \dots$ , given  $\mathbf{x}^{(t)} = (x^{(1)}, \dots, x^{(p)})$ , generate

1.  $X_1^{(t+1)} \sim f_1(x_1 | x_2^{(t)}, \dots, x_p^{(t)})$
2.  $X_2^{(t+1)} \sim f_2(x_2 | x_1^{(t+1)}, x_3^{(t)}, \dots, x_p^{(t)})$
- ...
- p.  $X_p^{(t+1)} \sim f_p(x_p | x_1^{(t+1)}, \dots, x_{p-1}^{(t+1)})$ .

The densities  $f_1, \dots, f_p$  are called the univariate conditional densities, and a particular feature of the Gibbs sampler is that these are the only densities used for simulation. Thus, even in a high-dimensional problem, all of the simulations can be univariate, which can be a huge advantage!

# Bayesian Change-Point Analysis


Let's try to model a more interesting example, a time series of recorded coal mining disasters in the UK from 1851 to 1962, for  $n = 111$  years of data. This analysis will follow the analysis layed out by Carlin et al in [Hierarchical Bayesian Analysis of Changepoint Problems](#). Occurrences of disasters in the time series is thought to be derived from a Poisson process with a drop rate parameter in the later part of the time series. We are interested in locating the change point,  $k$ , in the series.

## [Change Point Detection](#)

# Bayesian Change-Point Analysis

```
data_vector = c(4, 5, 4, 0, 1, 4, 3, 4, 0, 6, 3, 3, 4, 0, 2, 6,
               3, 3, 5, 4, 5, 3, 1, 4, 4, 1, 5, 5, 3, 4, 2, 5,
               2, 2, 3, 4, 2, 1, 3, 2, 2, 1, 1, 1, 1, 3, 0, 0,
               1, 0, 1, 1, 0, 0, 3, 1, 0, 3, 2, 2, 0, 1, 1, 1,
               0, 1, 0, 1, 0, 0, 0, 2, 1, 0, 0, 0, 1, 1, 0, 2,
               3, 3, 1, 1, 2, 1, 1, 1, 1, 2, 4, 2, 0, 0, 1, 4,
               0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1)
n = length(data_vector)
disasters = ts(data_vector, freq=1, start = 1851)
plot(disasters, type = "h", lty = 6, lwd = 2, main = "Coal Mining Disasters", xlab = "Year", ylab = "Disasters")
```

**Coal Mining Disasters**





# Bayesian Change-Point Analysis

We are going to use Poisson random variables for this type of count data. Denoting year  $i$ 's accident count by  $y_i$ ,

$$y_i \sim \text{Poisson}(\Lambda)$$

The modeling problem revolves around estimating the values of the  $\lambda$  parameters. Looking at the time series above, it appears that the rate declines later in the time series. A changepoint model identifies a point (year) during the observation period ( $k$ ) after which the parameter  $\lambda$  drops to a lower value. So we are estimating two  $\lambda$  parameters: one for the early period and another for the late period.

$$\Lambda = \begin{cases} \theta & \text{if } i < k \\ \lambda & \text{if } i \geq k \end{cases}$$

# Bayesian Change-Point Analysis

We need to assign prior probabilities to both  $\theta$  and  $\lambda$  parameters. The gamma distribution not only provides a continuous density function for positive numbers, but it is also conjugate with the Poisson sampling distribution.

We will specify suitably vague hyperparameters, letting  $\alpha = 1$  and allowing  $\beta$ s to vary.

$$\theta \sim \text{Gamma}(1, b_1)$$

$$\lambda \sim \text{Gamma}(1, b_2)$$

# Bayesian Change-Point Analysis

Since we do not have any intuition about the location of the changepoint (prior to viewing the data), we will assign a discrete uniform prior over all years 1851-1962.

$$k \sim \text{Unif}(1851, 1962)$$

$$\Rightarrow p(K = k) = \frac{1}{111}$$

# Implementing Gibbs sampling

We are interested in estimating the joint posterior of  $\theta, \lambda$  and  $k$  given the array of annual disaster counts  $\mathbf{y}$ . This gives:

$$p(\theta, \lambda, k | \mathbf{y}) \propto p(\mathbf{y} | \theta, \lambda, k) p(\theta, \lambda, k)$$

To employ Gibbs sampling, we need to factor the joint posterior into the product of conditional expressions:

$$p(\theta, \lambda, k | \mathbf{y}) \propto p(y_{i < k} | \theta, k) p(y_{i \geq k} | \lambda, k) p(\theta) p(\lambda) p(k)$$

which we have specified as:

$$\begin{aligned} p(\theta, \lambda, k | \mathbf{y}) &\propto \left[ \prod_{t=1851}^k \text{Poisson}(y_t | \theta) \times \prod_{t=k+1}^{1962} \text{Poisson}(y_t | \lambda) \right] \times \text{Gamma}(\theta | \alpha, \beta) \times \text{Gamma}(\lambda | \alpha, \beta) \frac{1}{111} \\ &\propto \left[ \prod_{t=1851}^k e^{-\theta} \theta^{y_t} \prod_{t=k+1}^{1962} e^{-\lambda} \lambda^{y_t} \right] \theta^{\alpha-1} e^{-\beta\theta} \lambda^{\alpha-1} e^{-\beta\lambda} \\ &\propto \theta^{\sum_{t=1851}^k y_t + \alpha - 1} e^{-(\beta+k)\theta} \lambda^{\sum_{t=k+1}^{1962} y_t + \alpha - 1} e^{-\beta\lambda} \end{aligned}$$

# Implementing Gibbs sampling

So, the full conditionals are known, and critically for Gibbs, can easily be sampled from.

$$\theta \sim \text{Gamma} \left( \sum_{t=1851}^k y_i + \alpha, k + \beta \right)$$

$$\lambda \sim \text{Gamma} \left( \sum_{t=k+1}^{1962} y_i + \alpha, 1962 - k + \beta \right)$$

$$p(k|\mathbf{y}, \theta, \lambda, b_1, b_2) = \frac{L(\mathbf{y}|\theta, \lambda, b_1, b_2)}{\sum_{j=1}^n L(\mathbf{y}|\theta, \lambda, b_1, b_2)}$$

where the likelihood is defined as

$$L(\mathbf{y}|\theta, \lambda, b_1, b_2) = e^{(\lambda - \theta)} \left( \frac{\theta}{\lambda} \right)^{\sum_1^k y_i}$$

# Implementing Gibbs sampling

```
set.seed(123)

y = data_vector
# Gibbs sampler for the coal mining change point
# initialization
n <- length(y) #length of the data
m <- 10^4 #length of the chain

## vectors to hold data
theta <- numeric(m)
lambda <- numeric(m)
k <- numeric(m)
L <- numeric(n)

## initial values
k[1] <- sample(1:n, 1) ## change-points
theta[1] <- 1
lambda[1] <- 1
a = 0.5
b1 <- 1
b2 <- 1
```

# Implementing Gibbs sampling

The algorithm explained by Carlin et al is simple. For  $t \in \{1, 2, \dots, m\}$

1. Sample  $\theta_t \sim \text{Gamma}(a_1 + \sum_1^k y_i, k_{t-1} + b_{1,t-1})$
2. Sample  $\lambda_t \sim \text{Gamma}(a_2 + \sum_{k+1}^n y_i, n - k_{t-1} + b_{2,t-1})$
3. Sample  $b_1 \sim \text{Gamma}(a_1 + c_1, (\theta_t + d_1))$
4. Sample  $b_2 \sim \text{Gamma}(a_2 + c_2, (\lambda_t + d_2))$
5. For  $j \in 1, \dots, n$  calculate  $L(\mathbf{y}|\theta, \lambda, b_1, b_2)$ , from there you'll obtain  $p(k|\mathbf{y}, \theta, \lambda, b_1, b_2)$
6. Sample  $k_t \sim p(k|\mathbf{y}, \theta, \lambda, b_1, b_2)$

# Implementing Gibbs sampling

```
# run the Gibbs sampler
for (t in 2:m){
  kt <- k[t-1]
  #generate theta
  r <- a + sum(y[1:kt])
  theta[t] <- rgamma(1, shape = r, rate = kt + b1)
  #generate lambda
  if (kt + 1 > n){
    r <- a + sum(y)
  }else{
    r <- a + sum(y[(kt+1):n])
  }
  lambda[t] <- rgamma(1, shape = r, rate = n - kt + b2)
  #generate b1 and b2
  b1 <- rgamma(1, shape = a, rate = theta[t]+1)
  b2 <- rgamma(1, shape = a, rate = lambda[t]+1)

  for (j in 1:n) {
    L[j] <- exp((lambda[t] - theta[t]) * j) *
      (theta[t] / lambda[t])^sum(y[1:j])
  }
  L <- L / sum(L)
  #generate k from discrete distribution L on 1:n
  k[t] <- sample(1:n, prob=L, size=1)
}
```



# Implementing Gibbs sampling

Set a burn-in of 1000 samples. We will use burn in to toss out "poor" samples from our Markov chain. Arguments for and against burn-in vary. Statisticians, Andrew Gelman ([Burn-in Man](#)) and Charlie Geyer ([Burn-In](#)) provide some commentary on burn-in.

```
## set burn in
burn_in <- 1000
## will toss out first 1000 samples
```

```
K <- k[burn_in:m]
```

```
## mean
print(mean(K))
```

```
## [1] 39.82791
```

```
#[1] 39.935
```

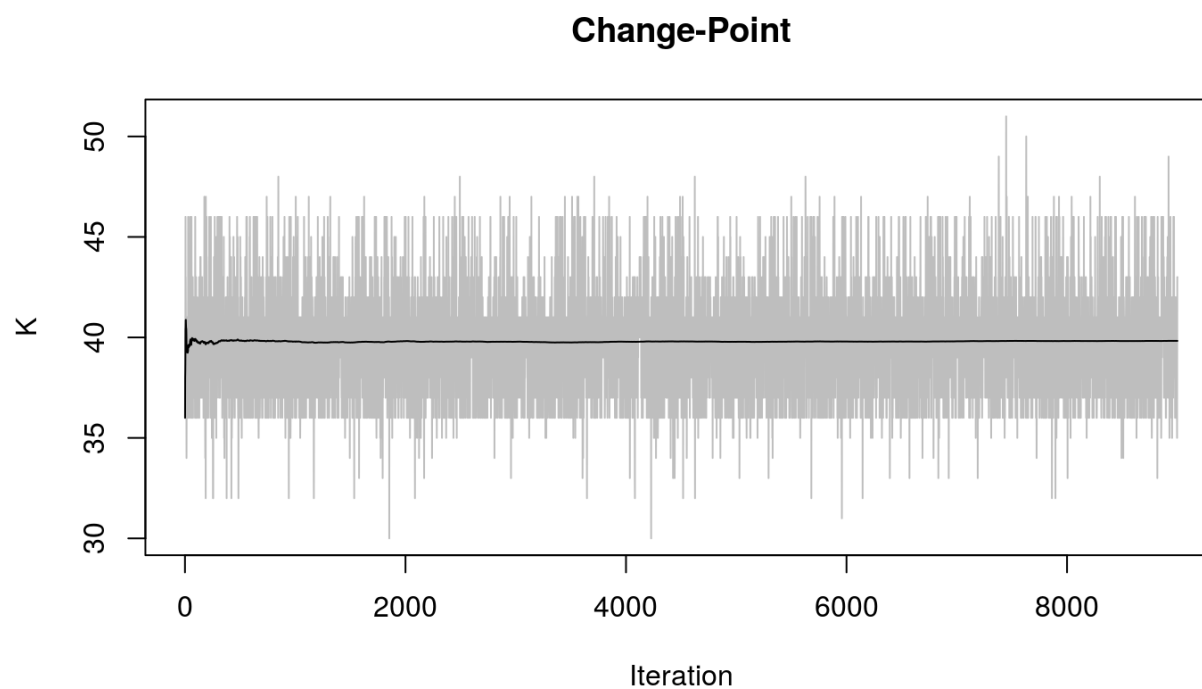
```
## mode
print(which.max(table(K)))
```

```
## 41
```

```
## 12
```

# Implementing Gibbs sampling

```
plot(K, type="l", col="gray", main = "Change-Point",  
     xlab="Iteration", ylab = "K")  
lines(1:length(K), cumsum(K) / (1:length(K)))
```



# Implementing Gibbs sampling

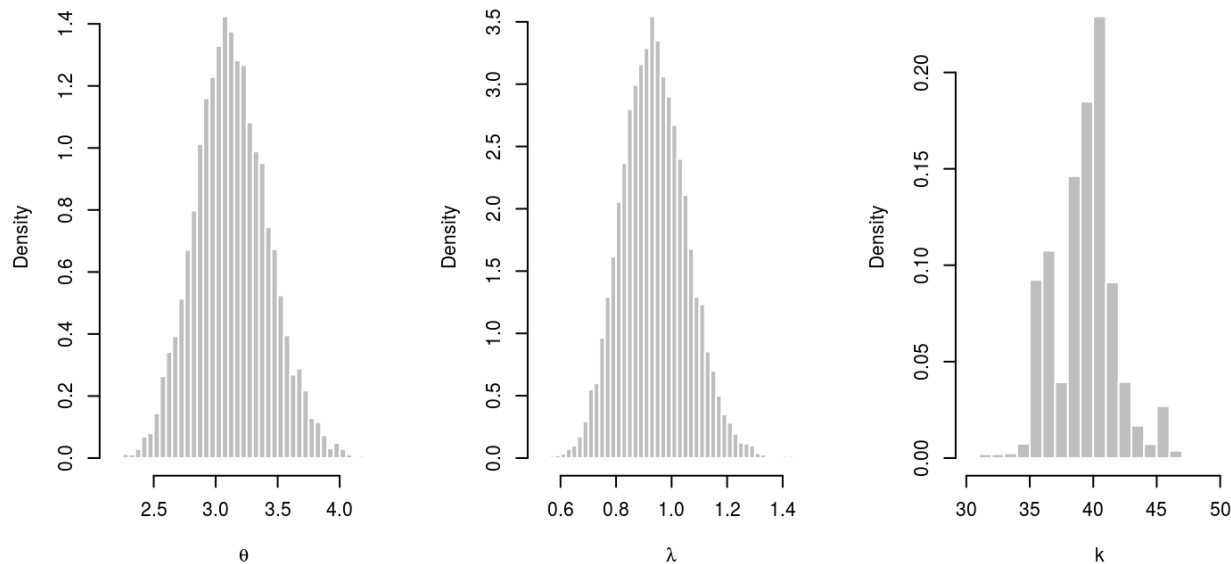
```
print(mean(lambda[burn_in:m])) #[1] 0.9341033
```

```
## [1] 0.939208
```

```
print(mean(theta[burn_in:m])) #[1] 3.108575
```

```
## [1] 3.130795
```

# Implementing Gibbs sampling



Our analysis can be used to justify that the change point occurs at some range at the 41<sup>st</sup> year, 1891, which is similar to other analyses cited by Carlin et al.

# Implementing Gibbs sampling

```
hist(theta[burn_in:m], main="", xlab = expression(mu),  
     col="gray", border="white",  
     breaks = "scott", prob=TRUE) #mu posterior  
  
hist(lambda[burn_in:m], main="", xlab = expression(lambda),  
     col="gray", border="white",  
     breaks = "scott", prob=TRUE) #lambda posterior  
  
hist(K, breaks=min(K):max(K), prob=TRUE, main="",  
     col="gray", border="white",  
     xlab = "k")
```