

Problem 1: This code generates 2000 sets of 5 variables, takes the sample mean of those 5 variables, and reports an estimate of the variance of the sample means at the end. Each row has the structure that $X_1 \sim N(0, 5)$, and $X_2, \dots, X_5 \sim N(0, 1)$. Therefore,

$$\begin{aligned}\text{var}\left(\frac{1}{n}(X_1 + \sum_{i=2}^5 X_i)\right) &= \frac{1}{n^2} \left(\text{var}(X_1) + \sum_{i=2}^5 \text{var}(X_i)\right) \\ &= \frac{1}{n^2} (5 + 1 + 1 + 1 + 1) \\ &= \frac{9}{25} = .36\end{aligned}$$

So entry entry of \mathbf{M} will be a $N(0, .36)$ random variable, and answer (4) will be the approximate output.

Problem 2: This program generates 1000 pairs of datasets of size $n = 20$. One of the datasets are iid $N(1, 1)$ and the other are iid $N(0, 1)$. A 95% confidence interval for $\mu_1 - \mu_2$ is formed for each of the 1000 datasets and I stores whether or not the confidence interval contained the true value for $\mu_1 - \mu_2$, which is 1 in this case. At the end the proportion of these confidence intervals that contained 1 is printed, which should be approximately .95.

Problem 3: First of all, this is an exactly quadratic function, so Newton-Raphson should converge to an extrema after a single iteration. Since this is a strictly non-positive function, 0 is clearly an extrema which is attained at $x = k$, which is where the algorithm should be after 1 iteration. This can be verified explicitly as well, since

$$f'(x) = -2(x - k)$$

and

$$f''(x) = -2$$

So, starting from x_0 , the next iteration takes you to

$$x_0 - \frac{2(x_0 - k)}{2} = x_0 - (x_0 - k) = k$$

Problem 4: Bisection begins by forming two new intervals based on splitting the bracket at the midpoint. So the two candidates for a new bracket are $(1/2, 3/4)$ and $(3/4, 1)$. From the information given in the problem there is not a sign change in f over the interval $(1/2, 3/4)$ and there is a sign change over the interval $(3/4, 1)$. Therefore, the new bracket will be $(3/4, 1)$.

Problem 5: Following from bilinearity of covariance,

$$\text{cov}(X + Y, X + Z) = \text{cov}(X, X) + \text{cov}(X, Z) + \text{cov}(Y, X) + \text{cov}(Y, Z) = \text{var}(X) = \lambda_1$$

which follows from the independence of X, Y, Z . A short R program to estimate this covariance, for particular values of λ_1, λ_2 , denoted by L1, L2 is

```
X <- rpois(1000, L1)
Y <- rpois(1000, L1)
Z <- rpois(1000, L2)
cov(X+Y, X+Z)
```

Problem 6: The code here estimates $\text{cov}(X, |X|)$, where $X \sim N(0, 1)$. The output should then be around $\text{cov}(X, |X|)$, which is

$$\text{cov}(X, |X|) = E(X \cdot |X|)$$

since X has mean 0. This can be calculated by using the law of total expectation, conditioning on whether or not X is positive:

$$\begin{aligned} E(X \cdot |X|) &= \frac{1}{2} (E(X \cdot |X| \mid X < 0) + E(X \cdot |X| \mid X > 0)) \\ &= \frac{1}{2} (-E(X^2) + E(X^2)) \\ &= 0 \end{aligned}$$

this is also intuitive because large values of $|X|$ can either be associated with very large values of X or very small values of X , so there couldn't possibly be any linear association between $|X|$ and X . So the approximate output should be 0. Notice this is an example of when uncorrelated does not imply independent.

Problem 7: For the first two blanks, you need to enter the R code to calculate the two estimators, which would be

```
1/mean(X)
```

and

```
sqrt(1/var(X))
```

respectively. The MSE is the average squared difference of the estimator of from the truth, so the second two blanks need to be filled in by the true value for λ , which is 1 in this case.

Problem 8: No, this does not mean that $\hat{\theta}_1$ will necessarily outperform $\hat{\theta}_2$ in terms of MSE, since

$$MSE(\hat{\theta}) = \text{bias}(\hat{\theta})^2 + \text{var}(\hat{\theta})$$

So, if $\hat{\theta}_2$ had substantially lower variance, then it could outperform $\hat{\theta}_1$ in terms of MSE.

Problem 9: We need to calculate the log-likelihood and its first two derivatives for this problem. The log-likelihood is

$$\ell(\theta) = \log(p_\theta(x)) = \text{const} + \frac{1}{2} \log(\theta) - \frac{\theta}{2x}$$

so the derivative of the log-likelihood (the score function) is

$$\ell'(\theta) = \frac{1}{2\theta} - \frac{1}{2x}$$

and the second derivative is

$$\ell''(\theta) = -\frac{1}{\theta^2}$$

So the first iteration of Newton Raphson will take you to

$$\begin{aligned}\theta_1 &= \theta_0 - \frac{\ell'(\theta_0)}{\ell''(\theta_0)} \\ &= \theta_0 + \frac{\frac{1}{2\theta_0} - \frac{1}{2x}}{-\frac{1}{\theta_0^2}} \\ &= \theta_0 + \frac{1}{2} - \frac{\theta_0}{2x} \\ &= .1 + \frac{1}{2} - \frac{.1}{2x} \\ &= .6 - \frac{.1}{2x}\end{aligned}$$

Problem 10: For the inversion method the first step is to calculate the CDF:

$$F(x) = \int_0^x \sin(y)dy = 1 - \cos(x)$$

It is straightforward to invert this function and see that

$$F^{-1}(y) = \cos^{-1}(1 - y)$$

so the R code to implement this is

```
rSin <- function(n) acos(1 - runif(n))
```

To do the rejection sampling optimally, we first need to determine

$$M = \max_{x \in (0, \pi/2)} \frac{p(x)}{g(x)} = \max_{x \in (0, \pi/2)} \frac{\pi \sin(x)}{2}$$

Since $\sin(x)$ is always at most 1, and $\sin(\pi/2) = 1$, this implies that

$$M = \pi/2$$

So $\frac{p(X)}{Mg(X)} = p(X) = \sin(X)$ an optimal rejection sampling algorithm will accept a candidate draw, $X \sim g$, when

$$U \leq \sin(X)$$

where U is Uniform(0,1). The optimal acceptance rate for rejection sampling is $1/M$, which in this case is $2/\pi \approx .636$.

Problem 11: This integral must evaluate to 0, because $\sinh(x)$ is an odd function— that is, $\sinh(x) = -\sinh(-x)$, and we are integrating over a region that is symmetric around 0, so the area to the left of 0 and the area to the right of 0 must cancel out.

This integral can be re-written as

$$I = 2\pi \int_{-\pi}^{\pi} \sinh(x)p(x)dx$$

where $p(x) = 1/2\pi$ is the $\text{Uniform}(-\pi, \pi)$ density. So,

$$I = 2\pi E(\sinh(U))$$

where $U \sim \text{Uniform}(-\pi, \pi)$. This expectation can be estimated by monte carlo by appealing to the law of large numbers. The following program does this:

```
U <- runif(1e4, -pi, pi)
2*pi*mean( sinh(U) )
```

Problem 12: Using the hint, if $X \sim p$, we can think of X in terms of the following algorithm:

1. Generate $U \sim \text{Uniform}(0, 1)$
2. If $U \leq .5$, then generate $X \sim N(-2, 1/2)$
3. If $U > .5$, generate $X \sim N(2, 1/2)$.

The following code does this for a pre-specified sample size, `n`:

```
U <- runif(n)
X <- rep(0,n)
w <- which(U <= .5)
X[w] <- rnorm(length(w), mean=-2, sd=sqrt(1/2))
X[-w] <- rnorm(n-length(w), mean=2, sd=sqrt(1/2))
```

Problem 13: Similarly to problem 11, this can be written as an expectation against the $\text{uniform}(0,1)$ distribution:

$$E(\sin(e^{1+U^2}))$$

Using $g(x; \alpha, \beta)$, the $\text{beta}(\alpha, \beta)$ density, as the trial density, define the importance function to be $w(x) = p(x)/g(x) = 1/g(x)$. Sample from $X \sim g$, and then estimate $E(w(X) \cdot \sin(e^{1+X^2}))$ by appealing to the law of large numbers. The following R program does this:

```
IS <- function(k, a, b)
{
  f <- function(x) sin(exp(1+x^2))
  p <- function(x) dunif(x)
  g <- function(x) dbeta(x,a,b)
  w <- function(x) p(x)/g(x)
  X <- rbeta(k, a, b)
  Q <- w(X)*f(X)
  return( c(mean(Q), var(Q)/k) )
}
```

We would choose α, β (**a, b** in the above program) by calculating $\text{IS}(\mathbf{k}, \mathbf{a}, \mathbf{b})$ for a fixed \mathbf{k} (say 10000), and comparing the standard error of the resulting estimate. The combination of α, β that corresponds to the smallest standard error would be considered the optimal choice.

Problem 14: Denoting the $N(0, \mu^2)$ density by $\phi_{\sigma^2}(x)$, this integral can be re-written as

$$\int_{-\infty}^{\infty} x^2 \frac{p(x)}{\phi_{\sigma^2}(x)} \phi_{\sigma^2}(x) dx$$

Letting $w(x) = \frac{p(x)}{\phi_{\sigma^2}(x)}$, this integral is the same as

$$E(X^2 w(X))$$

where $X \sim N(0, \sigma^2)$. This expectation can be estimated by appealing to the law of large numbers. The following R code carries out the importance sampling:

```
IS <- function(k, v)
{
  f <- function(x) x^2
  p <- function(x) exp(-x)/( (1+exp(-x))^2 )
  g <- function(x) dnorm(x, mean=0, sd=sqrt(v))
  w <- function(x) p(x)/g(x)
  X <- rnorm(k, mean=0, sd=sqrt(v))
  Q <- w(X)*f(X)
  return( c(mean(Q), var(Q)/k) )
}
```

To choose the optimal σ^2 , you would call `IS` for some fixed \mathbf{k} (say, 10000), and a grid of σ^2 values (say .01 up to 10, by .1), and see which value minimizes the variance of the integral approximation.

Problem 15: The partial derivatives of $f(x, y)$ are

$$\partial f / \partial x = (2x - 4y)f(x, y)$$

and

$$\partial f / \partial y = (2y - 4x)f(x, y)$$

The second derivatives are

$$\partial^2 f / \partial x^2 = (2 + (2x - 4y)^2)f(x, y),$$

$$\partial^2 f / \partial y^2 = (2 + (2y - 4x)^2)f(x, y),$$

and

$$\partial^2 f / \partial x \partial y = (-4 + (2x - 4y)(2y - 4x))f(x, y)$$

plugging into $x = 1, y = 1$ in accordance with the start values, the gradient at x_0 is

$$\nabla f(x_0) = (-2/e, -2/e)$$

and the hessian matrix at x_0 is

$$\nabla^2 f(x_0) = \begin{pmatrix} 6/e & 0 \\ 0 & 6/e \end{pmatrix}$$

The determinant of this matrix is $36/e^2$, so the inverse hessian is

$$(\nabla^2 f(x_0))^{-1} = \frac{e^2}{36} \begin{pmatrix} 6/e & 0 \\ 0 & 6/e \end{pmatrix} = \begin{pmatrix} e/6 & 0 \\ 0 & e/6 \end{pmatrix}$$

So the first iteration of Newton-Raphson will move you to

$$\begin{aligned} x_0 - (\nabla^2 f(x_0))^{-1} \nabla f(x_0) &= x_0 - \begin{pmatrix} e/6 & 0 \\ 0 & e/6 \end{pmatrix} (-2/e, -2/e) \\ &= (1, 1) + \begin{pmatrix} (2/e) \cdot (e/6) + 0 \cdot 2/e, 0 \cdot 2/e + (2/e) \cdot (e/6) \end{pmatrix} \\ &= (1, 1) + (1/3, 1/3) \\ &= (4/3, 4/3) \end{aligned}$$

Problem 16: Remember the posterior is equal to the likelihood multiplied by the prior. To begin, the likelihood is nothing but the joint density, so

$$\begin{aligned} L(\sigma^2) &= \prod_{i=1}^n p(x_i | \sigma^2) \\ &= \prod_{i=1}^n (2\pi\sigma^2)^{1/2} \cdot e^{-x_i^2/2\sigma^2} \\ &= (2\pi\sigma^2)^{n/2} \exp \left(-\frac{1}{2\sigma^2} \sum_{i=1}^n x_i^2 \right) \\ &\propto \sigma^n \exp \left(-\frac{1}{2\sigma^2} \sum_{i=1}^n x_i^2 \right) \end{aligned}$$

In this case the exponential density is the prior, so it is equal to

$$p(\sigma^2) = \lambda e^{-\lambda\sigma^2}$$

Therefore

$$\begin{aligned}
p(\sigma^2|x_1, \dots, x_n) &= L(\sigma^2) \cdot p(\sigma^2) \\
&= L(\sigma^2) \cdot \lambda e^{-\lambda\sigma^2} \\
&\propto \sigma^n \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^n x_i^2\right) \cdot \lambda e^{-\lambda\sigma^2} \\
&= \lambda \sigma^n \exp\left(-\lambda\sigma^2 - \frac{1}{2\sigma^2} \sum_{i=1}^n x_i^2\right)
\end{aligned}$$

is something proportional to the posterior distribution. Letting $g(\sigma^2)$ be the gamma(α, β) density and letting $\text{post}(\sigma^2)$ be the expression proportional to the posterior distribution, we would begin by determining

$$M = \max_{\sigma^2 > 0} \frac{\text{post}(\sigma^2)}{g(\sigma^2)}$$

Next you can generate a draw from $\text{post}(\sigma^2)$ by

1. Generate U from the Uniform(0,1) distribution
2. Generate Y from the gamma(α, β) distribution
3. If $U \leq \frac{\text{post}(Y)}{M \cdot g(Y)}$ then accept Y as a draw from $\text{post}(\sigma^2)$

Repeat this process until the desired sample size is achieved.